

cs120_lab2_linear_regression_df

databricks



(<http://creativecommons.org/licenses/by-nc-nd/4.0/>)

This work is licensed under a Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License.

(<http://creativecommons.org/licenses/by-nc-nd/4.0/>)

Workspace



Tables



Linear Regression Lab

This lab covers a common supervised learning pipeline, using a subset of the Million Song Dataset

(<http://labrosa.ee.columbia.edu/millionsong/>) from the UCI Machine Learning Repository

(<https://archive.ics.uci.edu/ml/datasets/YearPredictionMSD>). Our goal is to train a linear regression model to predict the release year of a song given a set of audio features.



This lab will cover:

- *Part 1:* Read and parse the initial dataset
 - *Visualization 1:* Features
 - *Visualization 2:* Shifting labels
- *Part 2:* Create and evaluate a baseline model

Send feedback