

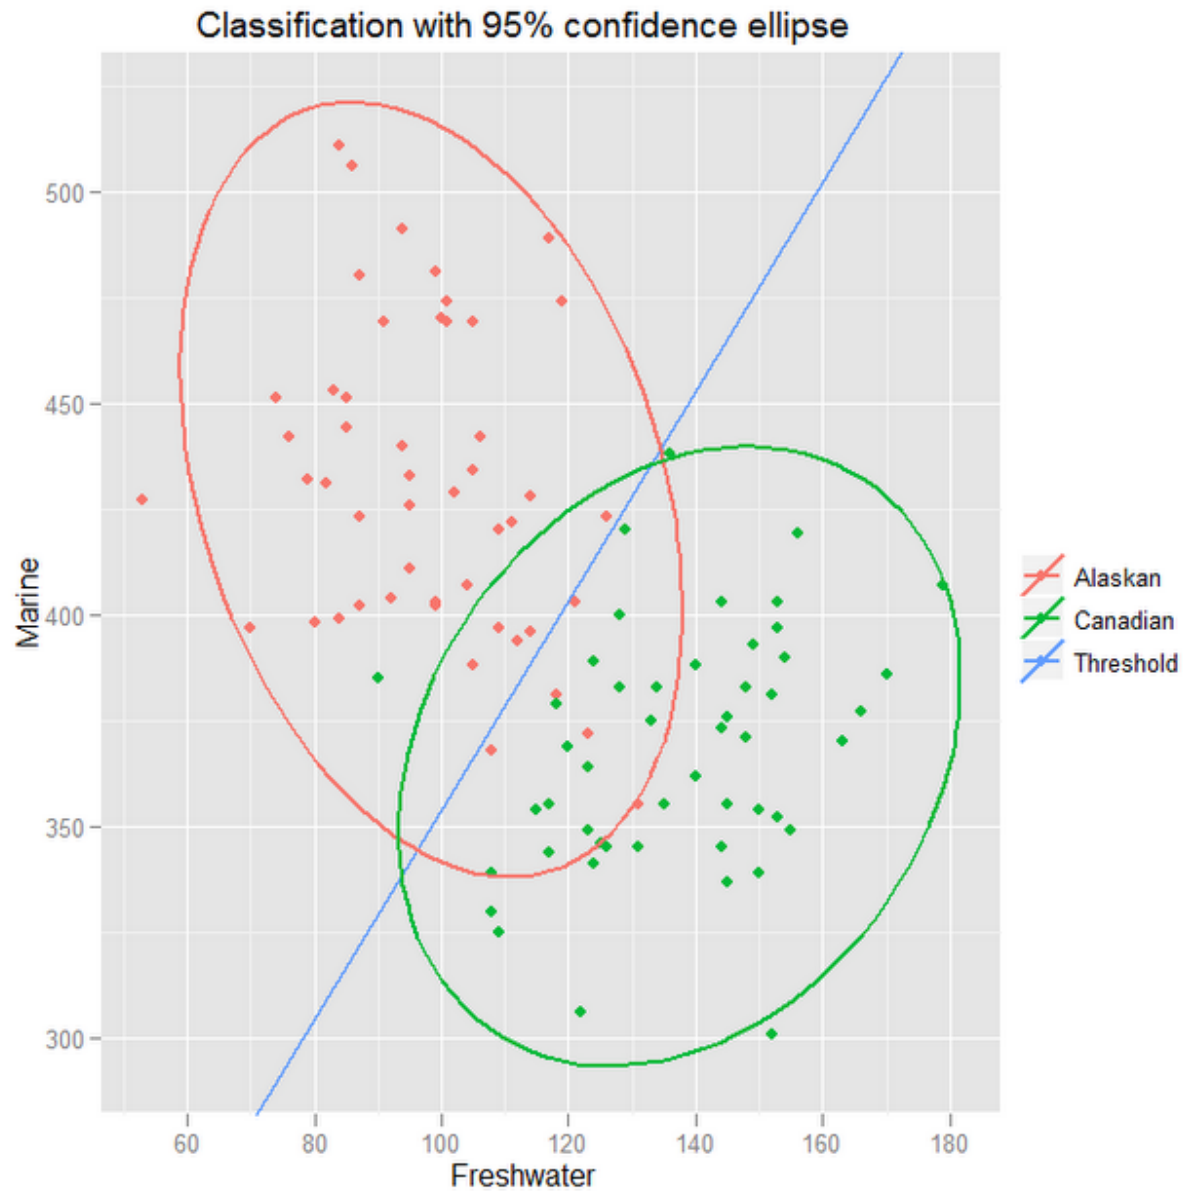


## multidimensional confidence intervals

---

I have numerous tuples (par1,par2), i.e. points in a 2 dimensional parameter space obtained from repeating an experiment multiple times.

I'm looking for a possibility to calculate and visualize confidence ellipses (not sure if thats the correct term for this). Here an example plot that I found in the web to show what I mean:



source: [blogspot.ch/2011/07/classification-and-discrimination-with.html](http://blogspot.ch/2011/07/classification-and-discrimination-with.html)

So in principle one has to fit a multivariate normal distribution to a 2D histogram of data points I guess. Can somebody help me with this?

python matplotlib scipy

asked Sep 6 '12 at 13:20



Raphael Roth

3,532 2 13 27

1 What's the input data? Is it an array of 2d points? Do you know in advance that there are 2 clusters? – [Daniel Velkov](#) Sep 6 '12 at 19:21

yes I know the number of clusters. I don't yet know what the format of the input data is, I guess a nx2 array where n is the number of points. – [Raphael Roth](#) Sep 7 '12 at 5:06

In that case you should cluster them first, then fit a gaussian to each cluster and finally plot the confidence intervals. Look at sklearn.cluster – [Daniel Velkov](#) Sep 7 '12 at 17:01

## 4 Answers

It sounds like you just want the 2-sigma ellipse of the scatter of points?

If so, consider something like this (From some code for a paper here:  
[https://github.com/joferkington/oost\\_paper\\_code/blob/master/error\\_ellipse.py](https://github.com/joferkington/oost_paper_code/blob/master/error_ellipse.py)):

```
import numpy as np

import matplotlib.pyplot as plt
from matplotlib.patches import Ellipse

def plot_point_cov(points, nstd=2, ax=None, **kwargs):
    """
    Plots an `nstd` sigma ellipse based on the mean and covariance of a point
    "cloud" (points, an Nx2 array).

    Parameters
    -----
    points : An Nx2 array of the data points.
    nstd : The radius of the ellipse in numbers of standard deviations.
           Defaults to 2 standard deviations.
    ax : The axis that the ellipse will be plotted on. Defaults to the
         current axis.
```

Additional keyword arguments are pass on to the ellipse patch.

#### Returns

-----

""" A matplotlib ellipse artist

"""

pos = points.mean(axis=0)

cov = np.cov(points, rowvar=False)

return plot\_cov\_ellipse(cov, pos, nstd, ax, \*\*kwargs)

def plot\_cov\_ellipse(cov, pos, nstd=2, ax=None, \*\*kwargs):

"""

Plots an `nstd` sigma error ellipse based on the specified covariance matrix (`cov`). Additional keyword arguments are passed on to the ellipse patch artist.

#### Parameters

-----

cov : The 2x2 covariance matrix to base the ellipse on

pos : The location of the center of the ellipse. Expects a 2-element sequence of [x0, y0].

nstd : The radius of the ellipse in numbers of standard deviations. Defaults to 2 standard deviations.

ax : The axis that the ellipse will be plotted on. Defaults to the current axis.

Additional keyword arguments are pass on to the ellipse patch.

#### Returns

-----

""" A matplotlib ellipse artist

"""

def eigsorted(cov):

vals, vecs = np.linalg.eigh(cov)

order = vals.argsort()[::-1]

return vals[order], vecs[:,order]

if ax is None:

ax = plt.gca()

vals, vecs = eigsorted(cov)

theta = np.degrees(np.arctan2(\*vecs[:,0][::-1]))

*# Width and height are "full" widths, not radius*


width, height = 2 \* nstd \* np.sqrt(vals)

ellip = Ellipse(xy=pos, width=width, height=height, angle=theta, \*\*kwargs)

ax.add\_artist(ellip)

return ellip

```
if __name__ == '__main__':  
    #-- Example usage -----  
    # Generate some random, correlated data  
    points = np.random.multivariate_normal(  
        mean=(1,1), cov=[[0.4, 9],[9, 10]], size=1000  
    )  
    # Plot the raw points...  
    x, y = points.T  
    plt.plot(x, y, 'ro')  
  
    # Plot a transparent 3 standard deviation covariance ellipse  
    plot_point_cov(points, nstd=3, alpha=0.5, color='green')  
  
    plt.show()
```



answered Sep 7 '12 at 15:39



Joe Kington

128k 19 299 307

nice, thanks for the answer. I hope I got this right: Assuming a multivariate normal distribution, one can simply take the eigenvalues and the eigenvectors to calculate the ellipses. – [Raphael Roth](#) Sep 10 '12 at 13:15

unfortunately, matplotlib patches cannot be drawn with logarithmic axes (or at least not correctly) as I need to .... why is life so complicated? – [Raphael Roth](#) Sep 10 '12 at 13:51

Yeah, I never thought to test it on logarithmic axes. One work-around would be to use a PathPatch, which will draw correctly on logarithmic axes. You'd have to generate points along the ellipse manually, but that's not too hard. – [Joe Kington](#) Sep 11 '12 at 0:15

@RaphaelRoth Another possibility to use logarithmic scale would be to fake it by transforming the datapoints and using a tick formatter for the axes (doesn't sound easy, but could be a way) – [heltonbiker](#) Oct 31 '12 at 13:02

- 2 @ThePredator - `arctan2` returns the full angle (can be in any of the 4 quadrants). `arctan` restricts the output to quadrants 1 and 4 (between  $-\pi/2$  and  $\pi/2$ ). You may notice that `arctan` takes a single parameter. Therefore, it can't distinguish between angles in quadrants 1 and 4 and a similar angle in quadrants 2 and 3. This is a convention that's shared by many other programming languages, in no small part because C defines them that way. – [Joe Kington](#) Jul 6 '15 at 11:44

|

I slightly modified one of the examples above that plots the error or confidence region contours. Now I think it gives the right contours.

It was giving the wrong contours because it was applying the `scoreatpercentile` method to the joint dataset (blue + red points) when it should be applied separately to each dataset.

The modified code can be found below:

```
import numpy
import scipy
import scipy.stats
import matplotlib.pyplot as plt

# generate two normally distributed 2d arrays
x1=numpy.random.multivariate_normal((100,420),[[120,80],[80,80]],400)
x2=numpy.random.multivariate_normal((140,340),[[90,-70],[-70,80]],400)
```

```

# fit a KDE to the data
pdf1=scipy.stats.kde.gaussian_kde(x1.T)
pdf2=scipy.stats.kde.gaussian_kde(x2.T)

# create a grid over which we can evaluate pdf
q,w=numpy.meshgrid(range(50,200,10), range(300,500,10))
r1=pdf1([q.flatten(),w.flatten()])
r2=pdf2([q.flatten(),w.flatten()])

# sample the pdf and find the value at the 95th percentile
s1=scipy.stats.scoreatpercentile(pdf1(pdf1.resample(1000)), 5)
s2=scipy.stats.scoreatpercentile(pdf2(pdf2.resample(1000)), 5)

# reshape back to 2d
r1.shape=(20,15)
r2.shape=(20,15)

# plot the contour at the 95th percentile
plt.contour(range(50,200,10), range(300,500,10), r1, [s1],colors='b')
plt.contour(range(50,200,10), range(300,500,10), r2, [s2],colors='r')

# scatter plot the two normal distributions
plt.scatter(x1[:,0],x1[:,1],alpha=0.3)
plt.scatter(x2[:,0],x2[:,1],c='r',alpha=0.3)

```

answered Jul 1 '13 at 16:29



Rodrigo

21 3

Refer the post [How to draw a covariance error ellipse](#).

Here's the python realization:

```

import numpy as np
from scipy.stats import norm, chi2

def cov_ellipse(cov, q=None, nsig=None, **kwargs):
    """
    Parameters
    -----
    cov : (2, 2) array
        Covariance matrix.
    q : float, optional

```

```

    Confidence level, should be in (0, 1)
    nsig : int, optional
        Confidence level in unit of standard deviations.
        E.g. 1 stands for 68.3% and 2 stands for 95.4%.

    Returns
    -----
    width, height, rotation :
        The lengths of two axes and the rotation angle in degree
        for the ellipse.
    """

    if q is not None:
        q = np.asarray(q)
    elif nsig is not None:
        q = 2 * norm.cdf(nsig) - 1
    else:
        raise ValueError('One of `q` and `nsig` should be specified.')
    r2 = chi2.ppf(q, 2)

    val, vec = np.linalg.eigh(cov)
    width, height = 2 * sqrt(val[:, None] * r2)
    rotation = np.degrees(arctan2(*vec[:, -1, 0]))

    return width, height, rotation

```

The meaning of *standard deviation* is **wrong** in the answer of Joe Kington. Usually we use 1, 2 sigma for 68%, 95% confidence levels, but the 2 sigma ellipse in his answer does not contain 95% probability of the total distribution. The correct way is using a chi square distribution to estimate the ellipse size as shown in the [post](#).

edited Sep 28 '16 at 14:09

answered Sep 28 '16 at 13:40



Syrtis Major

991 6 24

I guess what you are looking for is to compute the [Confidence Regions](#).

I don't know much how about it, but as a starting point, I would check the [sherpa](#) application for python. At least, in their Scipy 2011 talk, authors mention that you can determine and obtain confidence regions with it (you may need to have a model for your data though).

See the [video](#) and corresponding [slides](#) of the Sherpa talk.



HTH

answered Sep 6 '12 at 23:25

[gcalmettes](#)**4,106** 16 24

---

I also came along the sherpa-documentation, but I have actually no idea what this is :) – [Raphael Roth](#)  
Sep 7 '12 at 5:01

---