

Spatial, Temporal and Spatio-Temporal Data in R

L. Torgo

ltorgo@fc.up.pt

Faculdade de Ciências / LIAAD-INESC TEC, LA
Universidade do Porto

Jan, 2017



Jožef Stefan Institute

Introduction

Motivation

- We leave in a world where data collection devices abound
- Frequently these devices are location- and time-aware (e.g. smartphones)
- Data sets where observations are tagged by time and/or location are increasing at a fast rate



Jožef Stefan Institute

Motivation - 2

- If we suspect our observations may have some form of time and/or space correlation we should not neglect these effects in our analysis
 - Do observations of X at time t depend on the values at recent past times ?
 - Do observations of X at location z depend on values at neighboring locations ?
 - Do observations of X at location z on time t depend on the values at neighboring locations currently or in recent past times?

Some Tasks/Problems Addressed in Spatio-Temporal Data Mining

- Exploratory analysis of the data with the goal of understanding the eventual spatial and temporal patterns
- Prediction
 - Spatial interpolation
 - Time series forecasting
 - Spatio-temporal forecasting

Why do we need something specific?

- Spatial coordinates are just numbers...
- Temporal tags can be seen as strings following some rules...
- Special purpose classes facilitate the manipulation of this type of data
- e.g. temporal or spatial queries on a data set



Some Packages and Classes

- Temporal data
 - Package **xts**
- Spatial data
 - Package **sp**
- Spatio-temporal data
 - Package **spacetime**



Handling Time Series in R

- R includes several data structures that can be used to store time series
- In this illustration we will use the infra-structured provided in package `xts`
Note: this is an extra package that must be installed.
- The function `xts()` can be used to create a time series,

```
library(lubridate)
library(xts)
sp500 <- xts(c(1102.94, 1104.49, 1115.71, 1118.31),
              ymd(c("2010-02-25", "2010-02-26", "2010-03-01", "2010-03-02")))
sp500

##          [,1]
## 2010-02-25 1102.94
## 2010-02-26 1104.49
## 2010-03-01 1115.71
## 2010-03-02 1118.31
```

Creating time series

- The function `xts` has 2 arguments: the time series values and the temporal tags of these values
- The second argument must contain dates
- The function `ymd()` from package **lubridate** can be used to convert strings into dates
- If we supply a matrix on the first argument we will get a multivariate time series, with each column representing one of the variables

Indexing Time Series

- We may index the objects created by the function `xts()` as follows,

```
sp500[3]
##
[,1]
## 2010-03-01 1115.71
```

- However, it is far more interesting to make “temporal queries”,

```
sp500["2010-03-02"]
##
[,1]
## 2010-03-02 1118.31

sp500["2010-03"]
##
[,1]
## 2010-03-01 1115.71
## 2010-03-02 1118.31
```

Indexing Time Series (2)

```
sp500["2010-02-26/"]
##
[,1]
## 2010-02-26 1104.49
## 2010-03-01 1115.71
## 2010-03-02 1118.31

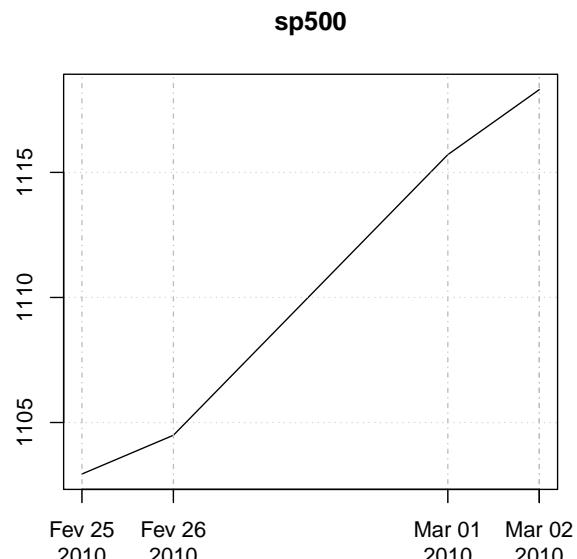
sp500["2010-02-26/2010-03-01"]
##
[,1]
## 2010-02-26 1104.49
## 2010-03-01 1115.71
```

- The index is a string that may represent intervals using the symbol `/` or by omitting part of a date. You may also use `::` instead of `/`.

Temporal Plots

- The `plot()` function can be used to obtain a temporal plot of a time series
- R takes care of selecting the proper axes,

```
plot(sp500)
```



Hands On Time Series

Package **quantmod** (an extra package that you need to install) contains several facilities to handle financial time series. Among them, the function `getMetals` allows you to download the prices of metals from `oanda.com`. Explore the help page of the function to try to understand how it works, and the answer the following:

- Obtain the prices of gold of the current year
- Show the prices in January
- Show the prices from February 10 till March 15
- Obtain the prices of silver in the last 30 days
Tip: explore the function `seq.Date()`
- Plot the prices of silver in the last 7 days
Tip: explore the function `last()` on package **xts**

Geo-referenced Data in R

- Package **sp** defines a series of classes of objects that can be used to store geo-referenced data sets
- Many other R packages build upon the classes defined in this package
- A broad overview of the (many) packages existing in R for this type of data can be found in
Analysis of Spatial Data Task View



Some of the Classes Defined in **sp**

- **SpatialPointsDataFrame**
A data frame like structure that can be used to store data (values of a set of variables) about geographic locations (a spatial point)
- **SpatialGridDataFrame**
A data frame like structure that can be used to store data (values of a set of variables) about a regular grid of spatial points
- **SpatialLinesDataFrame** and **SpatialPolygonsDataFrame**
A data frame like structure that can be used to store data (values of a set of variables) about lines and polygons (closed sequences of lines)



A Simple Illustration

Fires data of 500m² regions of Portugal

- Official data on fires in different regions of Portugal
- The data set we will use contains information on 25000 locations

```
df <- read.csv("firesnew_25000_500m.txt")
```

```
df[1:3,]

##   FID_ CID ano1991 ano1992 ano1993 ano1994 ano1995 ano1996 ano1997 ano1998
## 1    NA   1      0      0      0      0      0      0      0      0
## 2    NA   2      0      0      0      0      0      0      0      0
## 3    NA   3      0      0      0      0      0      0      0      0
##   ano1999 ano2000      x      y
## 1      0      0 -7.31924 38.5406
## 2      0      0 -7.63557 40.5022
## 3      0      0 -7.90273 40.3418
```



A Simple Illustration

Creating a `SpatialPointsDataFrame` object with 2000 data

```
library(sp)
spatialCoords <- cbind(long=df$x, lat=df$y)
coordRefSys <- CRS("+proj=longlat +ellps=WGS84")
fires2000 <- SpatialPointsDataFrame(spatialCoords,
                                     df[, "ano2000", drop=FALSE],
                                     proj4string=coordRefSys)
fires2000[1:3,]

##           coordinates ano2000
## 1 (-7.31924, 38.5406)      0
## 2 (-7.63557, 40.5022)      0
## 3 (-7.90273, 40.3418)      0
```



A Simple Illustration

Some utility functions

```
bbox(fires2000)

##           min      max
## long -9.49174 -6.20743
## lat  36.98050 42.14360

coordinates(fires2000)[1:3,]

##           long      lat
## [1,] -7.31924 38.5406
## [2,] -7.63557 40.5022
## [3,] -7.90273 40.3418

summary(fires2000)

## Object of class SpatialPointsDataFrame
## Coordinates:
##           min      max
## long -9.49174 -6.20743
## lat  36.98050 42.14360
## Is projected: FALSE
## proj4string : [+proj=longlat +ellps=WGS84]
## Number of points: 25000
## Data attributes:
##   ano2000
##   Min. :0.00000
##   Median :0.00000
##   Mean   :0.01612
##   3rd Qu.:0.00000
```

© L.Torgo (FCUP - LIAAD / UP)

Space-Time Data

Jan, 2017

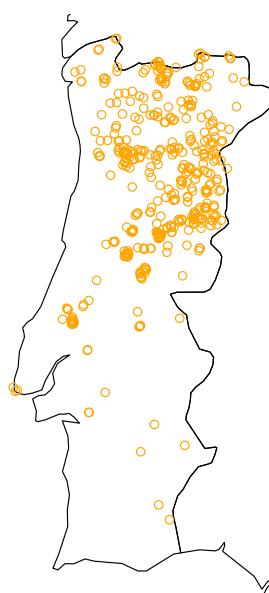
17 / 51

Geo-referenced Data

A Simple Illustration

Visualizing the data

```
library(rworldmap)
m <- getMap(resolution="low")
b <- bbox(fires2000)
plot(m, xlim=b[1,], ylim=b[2,])
points(fires2000[fires2000$ano2000==1,],
        col="orange")
title("Burnt Area during 2000 in Portugal")
```

Burnt Area during 2000 in Portugal

Spatio-temporal Data with **spacetime**

- Spatio-temporal data frequent formats
 - **Time-wide** - different columns have different measurements across time
 - **Space-wide** - different columns have measurements across different locations
 - **Long formats** - each data record contains measurements for a space-time combination
- Our previous fires data example used a time-wide format



Spatio-temporal Data with **spacetime** (cont.)

- Some types of spatio-temporal data sets
 - **Full grids** - we have values for all combinations of location and time stamps
number of values = number of locations \times number of time stamps
 - **Sparse grids** - similar but only non-missing values are stored
 - **Irregular data** - Each value is measured on a certain location and time, without regularity
 - **Time intervals, moving objects and trajectories** - e.g. spatial features constant but values collected on a time interval; trajectories where irregular space-time points form a sequence; etc.



Spatio-temporal Data with **spacetime** (cont.)

- The package **spacetime** defines different classes for these situations
 - **Full grids** - class STFDF
 - **Sparse grids** - class STSDF
 - **Irregular data** - class STIDF
 - **Time intervals, moving objects and trajectories** - class STTDF



A Simple Illustration with the Fires Data

The data again...

```
df <- read.csv("firesnew_25000_500m.txt")
```

```
df[1:3,]
```

```
##   FID_ CID ano1991 ano1992 ano1993 ano1994 ano1995 ano1996 ano1997 ano1998
## 1    NA   1      0      0      0      0      0      0      0      0      0
## 2    NA   2      0      0      0      0      0      0      0      0      0
## 3    NA   3      0      0      0      0      0      0      0      0      0
##   ano1999 ano2000     x     y
## 1      0      0 -7.31924 38.5406
## 2      0      0 -7.63557 40.5022
## 3      0      0 -7.90273 40.3418
```



Creating a STFDF object

- First let us put the data in Long Format

```
library(tidyrr)
x <- gather(df[,3:14], variable, value, ano1991:ano2000)
head(x, 3)

##           x      y variable value
## 1 -7.31924 38.5406  ano1991     0
## 2 -7.63557 40.5022  ano1991     0
## 3 -7.90273 40.3418  ano1991     0

x$variable <- substr(x$variable, 4, 7)
head(x, 3)

##           x      y variable value
## 1 -7.31924 38.5406      1991     0
## 2 -7.63557 40.5022      1991     0
## 3 -7.90273 40.3418      1991     0
```

Creating a STFDF object (cont.)

- We are now ready to create the object (we have a full grid)

```
library(sp)
library(spacetime)
spatialCoords <- cbind(long=df$x, lat=df$y)
coordRefSys <- CRS("+proj=longlat +ellps=WGS84")
timeStamps <- as.POSIXct(unique(x$variable),
                         format="%Y", tz="GMT")
sp <- SpatialPoints(spatialCoords, coordRefSys)
## STFDF assumes that space "moves faster" than time in the data
std <- STFDF(sp, timeStamps, data=x[, "value", drop=FALSE])
```

Querying the spatio-temporal data set

■ Data for a certain location

```
std[1,]

##           value timeIndex
## 1991-01-12     0        1
## 1992-01-12     0        2
## 1993-01-12     0        3
## 1994-01-12     0        4
## 1995-01-12     0        5
## 1996-01-12     0        6
## 1997-01-12     0        7
## 1998-01-12     0        8
## 1999-01-12     0        9
## 2000-01-12     0       10
```



Querying the spatio-temporal data set

■ Data for all locations on a certain time stamp

```
dataSince1995 <- std[, "1995/"]
data1991 <- std[, "1991"]
data1991[1:4,]

##           coordinates value
## 1 (-7.31924, 38.5406) 0
## 2 (-7.63557, 40.5022) 0
## 3 (-7.90273, 40.3418) 0
## 4 (-7.25657, 39.2572) 0
```

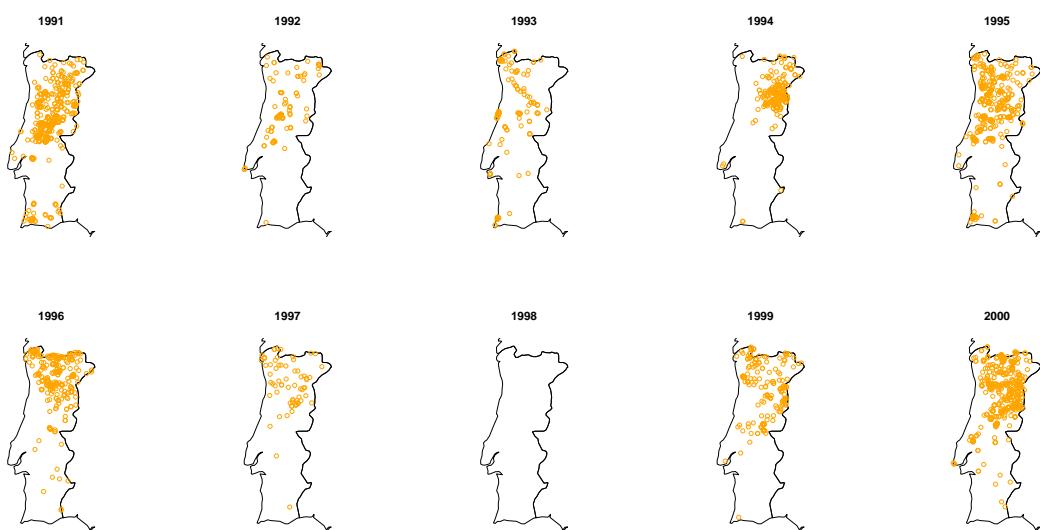


Plotting the spatio-temporal data set

```
library(rworldmap)
m <- getMap(resolution="low")
par(mfrow=c(2,5))
for(year in 1991:2000) {
  y <- std[,as.character(year)]
  bb <- bbox(y)
  plot(m,xlim=bb[1,],ylim=bb[2,])
  points(y[y$value ==1,],
          col="orange")
  title(year)
}
par(mfrow=c(1,1))
```



Plotting the spatio-temporal data set (cont.)



The Package **ggmap**

- Spatial visualization is a key step on spatio-temporal data analysis
- **ggmap** is a recent package that greatly enhances spatial visualization in R
- The package is able to combine spatial information from providers of maps (e.g. Google Maps) with the beautiful graphics of the **ggplot2** package
- On top of spatial visualization tools the package also includes also other nice utility functions



What is **ggplot2**

- An alternative graphics system for R based on the notion of *layered grammar of graphics*
- Every plot is formed by five components:
 - A data set and a set of aesthetic mappings (the way variables in the data are mapped into things we see in the graph)
 - One or more layers, each with a geometric object (*geom*) and a statistical transformation (*stat*)
 - A scale for each aesthetic mapping (frequently automatically generated)
 - A coordinate systems
 - A facet specification



ggmap and ggplot2

- As **ggmap** uses **ggplot2** graphics, the graphs also have the same components
- However, some of these are fixed to map components:
 - The x aesthetic is fixed to longitude
 - The y aesthetic is fixed to latitude
 - The coordinate system is fixed to the Mercator projection



The process of using **ggmap**

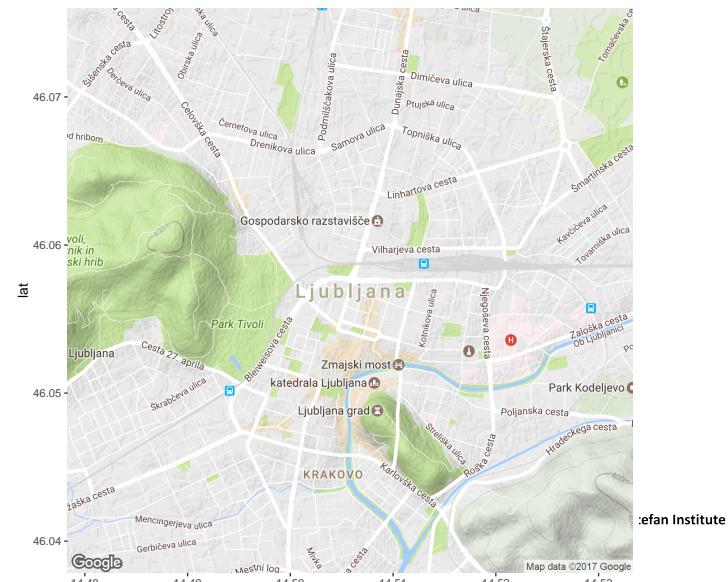
- Download a map image
- Plot it as basic layer using **ggplot2**
- Plot additional layers of data, statistics or models on top
- In **ggmap** this is done using:
 - `get_map()` to obtain the map image
 - `ggmap()` to make the actual plot



The function get_map()

- **ggmap** can use several map providers
Google Maps, OpenStreet Maps, Stamen Maps and CloudMade Maps
- A simple example

```
library(ggmap)
map <- get_map("Ljubljana",
                zoom=14)
ggmap(map)
```



© L.Torgo (FCUP - LIAAD / UP)

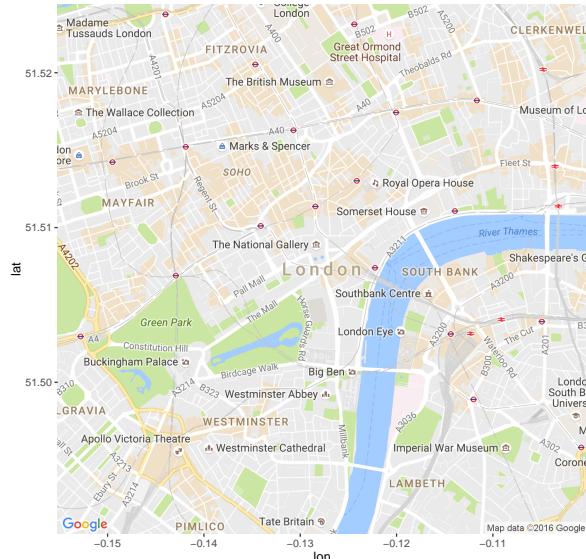
Space-Time Data

Jan, 2017

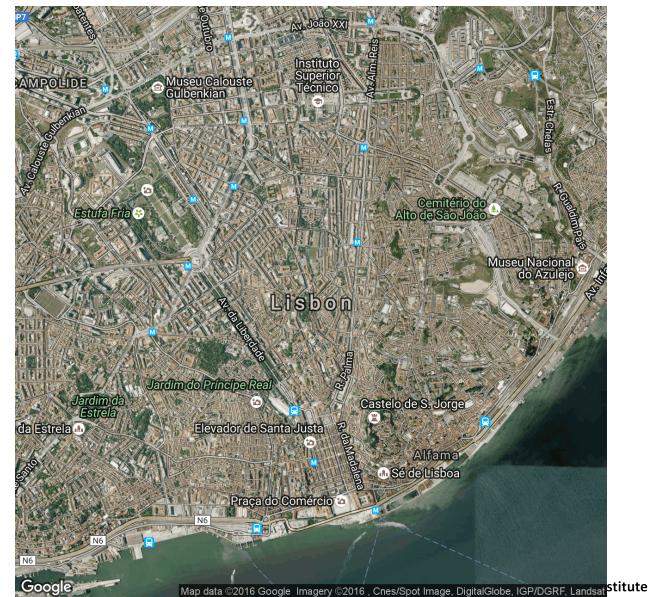
33 / 51

The function get_map() - a few examples

```
library(ggmap)
map <- get_map("London", zoom=14,
                maptype="roadmap")
ggmap(map)
```



```
map <- get_map("Lisbon", zoom=14,
                maptype="hybrid")
ggmap(map, extent="device")
```



© L.Torgo (FCUP - LIAAD / UP)

Space-Time Data

Jan, 2017

34 / 51

A Simple Illustration

Fires data of 500m² regions of Portugal

- Official data on fires in different regions of Portugal
- The data set we will use contains information on 25000 locations

```
df <- read.csv("firesnew_25000_500m.txt")
```

```
df[1:3, ]
```

```
##      FID_ CID ano1991 ano1992 ano1993 ano1994 ano1995 ano1996 ano1997 ano1998
## 1     NA    1       0       0       0       0       0       0       0       0
## 2     NA    2       0       0       0       0       0       0       0       0
## 3     NA    3       0       0       0       0       0       0       0       0
##   ano1999 ano2000      x      y
## 1       0       0 -7.31924 38.5406
## 2       0       0 -7.63557 40.5022
## 3       0       0 -7.90273 40.3418
```



Putting the data in long format

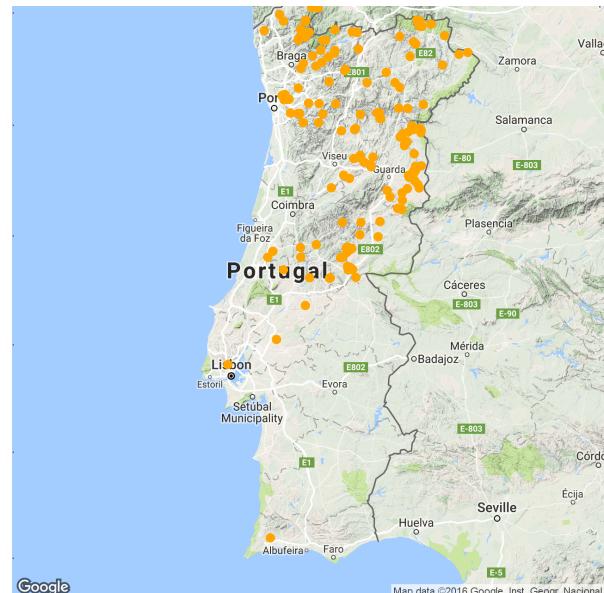
```
library(tidyr)
dat <- gather(df[, 3:14], year, burnt, ano1991:ano2000)
dat$year <- substr(dat$year, 4, 7)
head(dat, 4)
```

```
##           x      y year burnt
## 1 -7.31924 38.5406 1991     0
## 2 -7.63557 40.5022 1991     0
## 3 -7.90273 40.3418 1991     0
## 4 -7.25657 39.2572 1991     0
```



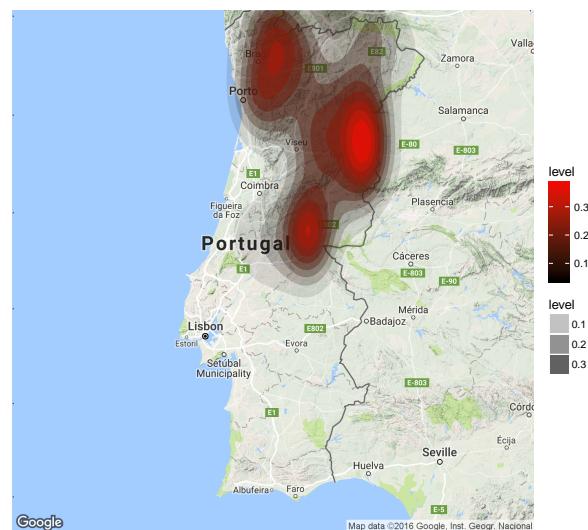
Plotting the Fires in 1999

```
library(ggmap)
pt <- get_map("Portugal", zoom=7)
data2plot <- dat[dat$year==1999 & dat$burnt == 1, ]
ggmap(pt, extent="device") +
  geom_point(data=data2plot,
             aes(x=x, y=y),
             color="orange", size=3)
```



Adding some spatial interpolation

```
ggmap(pt, extent="device",
      base_layer=ggplot(data2plot, aes(x=x, y=y)) +
      stat_density2d(data=data2plot,
                     aes(x=x, y=y,
                         fill=..level..,
                         alpha=..level..),
                     bins=10,
                     geom="polygon") +
      scale_fill_gradient(low="black", high="red"))
```



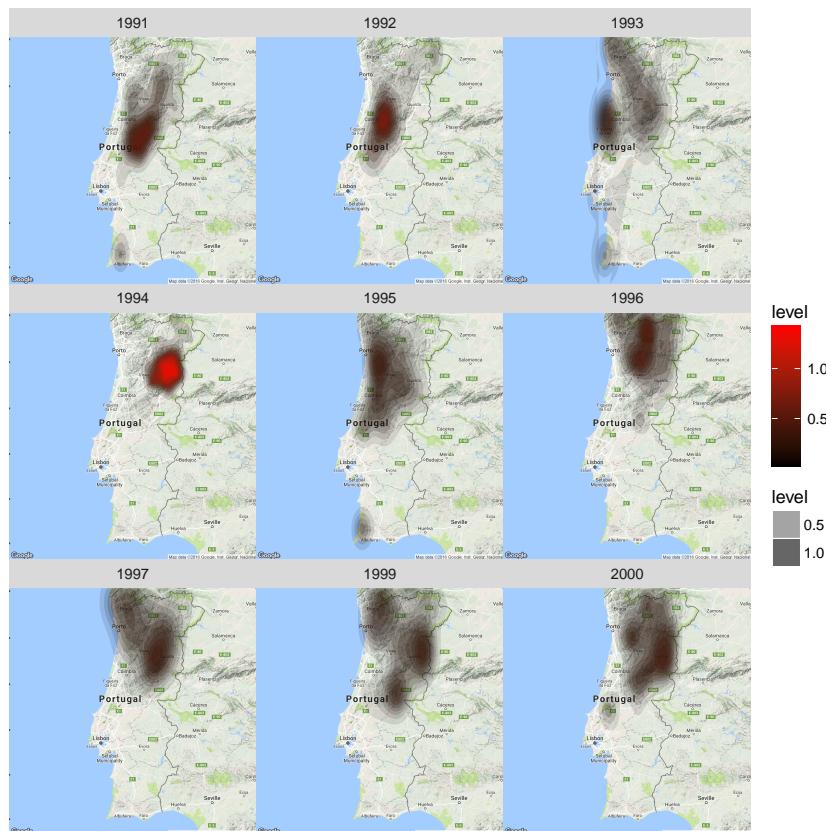
Spatio-temporal Visualization

- We need to add facetting along the year an everything else stays the same

```
data2plot <- dat[dat$burnt == 1, ]
ggmap(pt, extent="device",
      base_layer=ggplot(data2plot, aes(x=x, y=y))
    ) +
  stat_density2d(data=data2plot,
                 aes(x=x, y=y,
                     fill=..level..,
                     alpha=..level..),
                 bins=10,
                 geom="polygon") +
  scale_fill_gradient(low="black", high="red") +
  facet_wrap(~ year)
```



Spatio-temporal Visualization (cont.)



Getting geographic information of a location

- Taking advantage of the Google Maps API we can obtain information on a location

```
geocode("Ljubljana")

##      lon      lat
## 1 14.50575 46.05695

geocode("Ljubljana", output="more")

##      lon      lat   type   loctype      address      north
## 1 14.50575 46.05695 locality approximate ljubljana, slovenia 46.1422
##      south    east    west locality administrative_area_level_1 country
## 1 45.99007 14.6447 14.4195 Ljubljana                 Ljubljana Slovenia
```

Note: Please note that the Google Maps API has several request limitations. Namely, it has an unspecified short-term rate limit as well as a 24-hour limit of 2500 requests.



Getting physical addresses from a long/lat pair

- Checking the address of one of the fire locations

```
revgeocode(as.numeric(df[1000, c("x", "y")]))

## [1] "Unnamed Road, 7370, Portugal"
```



Calculating distances between locations

- The distance from New York to three other cities
- Note that the default mode is driving
- Again several request limitations are imposed by Google

```
from <- rep("New York", 3)
to <- c("Los Angeles", "San Francisco", "Toronto")
mapdist(from, to)

##      from          to     m      km    miles seconds minutes
## 1 New York    Los Angeles 4493003 4493.003 2791.9521 144653 2410.883
## 2 New York San Francisco 4674274 4674.274 2904.5939 152102 2535.033
## 3 New York        Toronto  790159   790.159  491.0048  27837  463.950
##      hours
## 1 40.18139
## 2 42.25056
## 3  7.73250

mapdist("Chinatown", "Times Square", mode="walking")

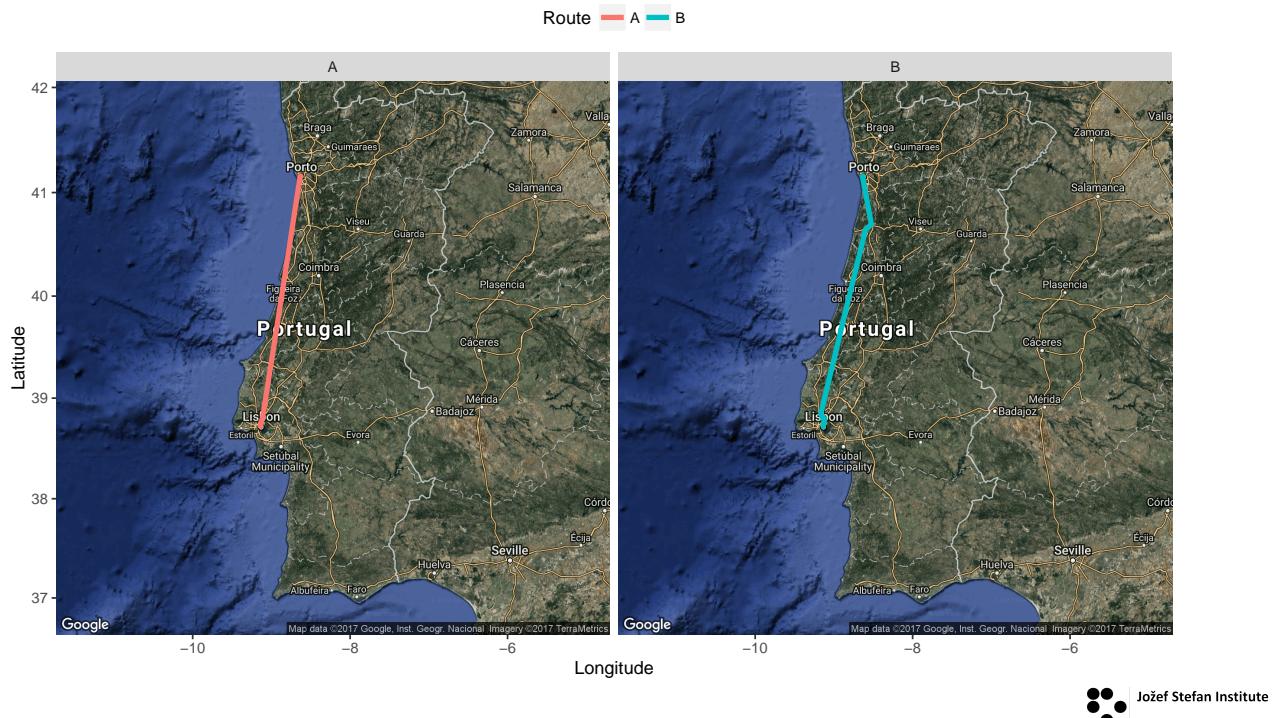
##      from          to     m      km    miles seconds minutes hours
## 1 Chinatown Times Square 5206 5.206 3.235008     3903   65.05 1.084167
```

Routes between locations

- Note that the default mode is driving
- Again several request limitations are imposed by Google

```
rtDat <- route("Porto", "Lisbon", mode="driving", structure="route", alternatives=TRUE)
mp <- get_map("Portugal", zoom=7, maptype="hybrid")
ggmap(mp,
  base_layer=ggplot(rtDat, aes(x=lon, y=lat, colour=route))) +
  geom_path(size=1.5, lineend="round") +
  facet_wrap(~ route) +
  labs(x = "Longitude", y = "Latitude", colour = "Route") +
  theme(legend.position = "top")
```

Routes between locations (cont.)



Interactive spatial visualization using package leaflet

- Leaflet is a very popular open-source JavaScript library for interactive maps
- The R package **leaflet** allows you to create this type of graphs in R
- After installing the package we can try it using the forest fires data

Interactive spatial visualization of the forest fires

```
df <- read.csv("firesnew_25000_500m.txt")
```

- The following shows the fires in 1999 in an interactive map

```
library(sp)
spatialCoords <- cbind(long=df$x, lat=df$y)
coordRefSys <- CRS("+proj=longlat +ellps=WGS84")
fires1999 <- SpatialPointsDataFrame(spatialCoords,
                                     df[, "ano1999", drop=FALSE],
                                     proj4string=coordRefSys)

library(leaflet)
leaflet() %>%
  addTiles() %>%
  addCircleMarkers(data=fires1999[fires1999$ano1999==1, ])
```



Improving a bit the visualization

- Getting the addresses of the places where there were fires

```
library(ggmap)
placesWithFires <- which(fires1999$ano1999 == 1)
coord <- coordinates(fires1999)[placesWithFires, ]
## Note that the Google API imposes limits on the following...
adds <- apply(coord, 1, function(cs) revgeocode(cs)))
## Now the interactive map (try clicking on a dot)
leaflet() %>%
  addTiles() %>%
  addCircleMarkers(data=fires1999[fires1999$ano1999==1, ],
                   popup=adds,
                   clusterOptions = markerClusterOptions())
```



Learning more about leaflet

You may learn more at <https://rstudio.github.io/leaflet/>



Hands On ggmap

Hands On Spatio-Temporal Data with ggmap

The file `irishWind.Rdata` contains two data frames with information on wind data values collected in several meteorological stations in Ireland along several years. The data frame **wind** contains the wind values for the different stations (in wide format), while the **wind.loc** data frame contains information on the stations. Using this data set answer the following questions:

- 1** Obtain the geographic coordinates of the stations
- 2** Reproduce the graph to the right



Hands On Spatio-Temporal Data with ggmap (cont.)

- 3 Using the functionalities provided by packages **tidyr** and **dplyr** obtain a data frame with the average yearly wind speed for each station.
- 4 Produce a spatio-temporal showing theses yearly averages on the stations.

