

## Announcing Stack Overflow Documentation

We started with Q&A. Technical documentation is next, and we need your help.

Whether you're a beginner or an experienced developer, you *can* contribute.

[Sign up and start helping →](#)[Learn more about Documentation →](#)

## Difference between map, applymap and apply methods in Pandas

Can you tell me when to use these vectorization methods with basic examples? I see that `map` is a `Series` method whereas the rest are `DataFrame` methods. I got confused about `apply` and `applymap` methods though. Why do we have two methods for applying a function to a `DataFrame`? Again, simple examples which illustrate the usage would be great!

Thanks!

[python](#) [numpy](#) [pandas](#) [vectorization](#)

edited Jul 16 '14 at 18:18

asked Nov 5 '13 at 20:20



[marillion](#)

1,276 6 19 30

### 5 Answers

Straight from Wes McKinney's [Python for Data Analysis](#) book, pg. 132 (I highly recommended this book):

Another frequent operation is applying a function on 1D arrays to each column or row. DataFrame's apply method does exactly this:

```
In [116]: frame = DataFrame(np.random.randn(4, 3), columns=list('bde'), index=['Utah', 'Ohio', 'Texas', 'Oregon'])
```

```
In [117]: frame
```

```
Out[117]:
```

	b	d	e
Utah	-0.029638	1.081563	1.280300
Ohio	0.647747	0.831136	-1.549481
Texas	0.513416	-0.884417	0.195343
Oregon	-0.485454	-0.477388	-0.309548

```
In [118]: f = lambda x: x.max() - x.min()
```

```
In [119]: frame.apply(f)
```

```
Out[119]:
```

	b	d	e
b	1.133201		
d	1.965980		
e	2.829781		

dtype: float64

Many of the most common array statistics (like sum and mean) are DataFrame methods, so using apply is not necessary.

Element-wise Python functions can be used, too. Suppose you wanted to compute a formatted string from each floating point value in frame. You can do this with applymap:

```
In [120]: format = lambda x: '%.2f' % x
```

```
In [121]: frame.applymap(format)
```

```
Out[121]:
```

	b	d	e
Utah	-0.03	1.08	1.28
Ohio	0.65	0.83	-1.55
Texas	0.51	-0.88	0.20
Oregon	-0.49	-0.48	-0.31

The reason for the name applymap is that Series has a map method for applying an element-wise function:

```
In [122]: frame['e'].map(format)
Out[122]:
Utah      1.28
Ohio     -1.55
Texas      0.20
Oregon    -0.31
Name: e, dtype: object
```

Summing up, `apply` works on a row / column basis of a DataFrame, `applymap` works element-wise on a DataFrame, and `map` works element-wise on a Series.

answered Nov 5 '13 at 20:40



[jeremiahbuddha](#)


2,124 1 13 22

- 
- 5 strictly speaking, `applymap` internally is implemented via `apply` with a little wrap-up over passed function parameter (roughly speaking replacing `func` to `lambda x: [func(y) for y in x]`, and applying column-wise) – [alko](#) Nov 5 '13 at 20:53
- 
- 1 Thanks for the explanation. Since `map` and `applymap` both work element-wise, I would expect a single method (either `map` or `applymap`) which would work both for a Series and a DataFrame. Probably there are other design considerations, and Wes McKinney decided to come up with two different methods. – [marillion](#) Nov 5 '13 at 21:58
- 

It's on page 129 in my copy for some reason. There's no label for second edition or anything. – [Jody](#) Jan 26 at 21:34

---

Work on work you love. From home.



@jeremiahbuddha mentioned that `apply` works on row/columns, while `applymap` works element-wise. But it seems you can still use `apply` for element-wise computation....

```
frame.apply(np.sqrt)
Out[102]:
```

	b	d	e
Utah	NaN	1.435159	NaN
Ohio	1.098164	0.510594	0.729748

```
Texas      NaN    0.456436    0.697337
Oregon    0.359079      NaN      NaN
```

```
frame.applymap(np.sqrt)
```

```
Out[103]:
```

```
      b      d      e
Utah   NaN  1.435159   NaN
Ohio  1.098164  0.510594  0.729748
Texas   NaN  0.456436  0.697337
Oregon  0.359079   NaN   NaN
```

answered Dec 19 '13 at 17:21



[user2921752](#)

166 6

- 7 Good catch with this. The reason this works in your example is because `np.sqrt` is a ufunc, i.e. if you give it an array, it will broadcast the `sqrt` function onto each element of the array. So when `apply` pushes `np.sqrt` on each columns, `np.sqrt` works itself on each of the elements of the columns, so you are essentially getting the same result as `applymap`. – [jeremiahbuddha](#) Jan 16 '14 at 0:22

Adding to the other answers, in a `Series` there are also [map](#) and [apply](#).

**Apply can make a DataFrame out of a series;** however, `map` will just put a series in every cell of another series, which is probably not what you want.

```
In [41]: p=pd.Series([1,2,3])
```

```
In [42]: p.apply(lambda x: pd.Series([x, x]))
```

```
Out[42]:
```

```
0 1
0 1 1
1 2 2
2 3 3
```

```
In [43]: p.map(lambda x: pd.Series([x, x]))
```

```
Out[43]:
```

```
0 0 1
1 1
dtype: int64
1 0 2
1 2
dtype: int64
```

```
2    0    3
1    3
dtype: int64
dtype: object
```

Also if I had a function with side effects, such as "connect to a web server", I'd probably use `apply` just for the sake of clarity.

```
series.apply(download_file_for_every_element)
```

Map **can use not only a function, but also a dictionary or another series**. Let's say you want to manipulate [permutations](#).

Take

```
1 2 3 4 5
2 1 4 5 3
```

The square of this permutation is

```
1 2 3 4 5
1 2 5 3 4
```

You can compute it using `map`. Not sure if self-application is documented, but it works in `0.15.1`.

```
In [39]: p=pd.Series([1,0,3,4,2])
```

```
In [40]: p.map(p)
```

```
Out[40]:
0    0
1    1
2    4
3    2
4    3
dtype: int64
```

answered Dec 8 '14 at 23:30



osa

2,677 19 28

Just wanted to point out, as I struggled with this for a bit

```
def f(x):  
    if x < 0:  
        x = 0  
    elif x > 100000:  
        x = 100000  
    return x
```

```
df.applymap(f)  
df.describe()
```

**this does not modify the dataframe itself, has to be reassigned**

```
df = df.applymap(f)  
df.describe()
```

answered Sep 26 '15 at 1:30



muon

168 10

---

I sometimes have trouble in figuring out whether you have to reassign or not after doing something with the df. It's mostly trial and error for me, but I bet there is a logic to how it works (that I am missing out). –

[marillion](#) Apr 13 at 16:19

---

in general, a pandas dataframe is only modified by either reassigning `df = modified_df` or if you set `inplace=True` flag. Also dataframe will change if you pass a dataframe to a function by reference and the function modifies the dataframe – [muon](#) Apr 13 at 16:30

---

This is not entirely true, think of `.ix` or `.where` etc. Not sure what the full explanation is for when you need to re-assign and when not. – [Thanos](#) Apr 14 at 20:45

---

Probably simplest explanation the difference between `apply` and `applymap`:

**`apply`** takes the whole column as a parameter and then assign the result to this column

**`applymap`** takes the separate cell value as a parameter and assign the result back to this cell.

NB If `apply` returns the single value you will have this value instead of the column after assigning and eventually will have just a row instead of matrix.

answered May 20 at 2:10



Kath

1,379 7 13