

Source code for cntk.learners

```
# Copyright (c) Microsoft. All rights reserved.
# Licensed under the MIT license. See LICENSE.md file in the project root
# for full license information.
# =====
```

```
'''
```

A learner tunes a set of parameters during the training process. One can use different learners for different sets of parameters. Currently, CNTK supports the following learning algorithms:

```
- :func:`AdaDelta <adadelat>`
- :func:`AdaGrad <adagrad>`
- :func:`FSAdaGrad <fsadagrad>`
- :func:`Adam <adam>`
- :func:`MomentumSGD <momentum_sgd>`
- :func:`Nesterov <nesterov>`
- :func:`RMSProp <rmsprop>`
- :func:`SGD <sgd>`
- :func:`Learner with a customized update function <universal>`
'''
```

```
from enum import Enum, unique
import warnings
import numpy as np
import cntk.internal.utils as utils

from .. import cntk_py, NDArrayView, asarray
from cntk.internal import typemap
from ..internal.swig_helper import map_if_possible
```

```
@unique
class UnitType(Enum):
```

[\[docs\]](#)

```
'''
```

Indicates whether the values in the schedule are specified on the per-sample or per-minibatch basis.

```
'''
```

```
sample = 'sample'
'''
```

Schedule contains per-sample values.

```
'''
```

```
minibatch = 'minibatch'
'''
```

Schedule contains per-minibatch values (and need to be re-scaled by the learner using the actual minibatch size in samples).

```
'''
```

```
def default_unit_gain_value():
```

[\[docs\]](#)

```
'''
```

Returns true if by default momentum is applied in the unit-gain fashion.

```
'''
```

```
return cntk_py.default_unit_gain_value()
```

```
def set_default_unit_gain_value(value):
```

[\[docs\]](#)

```
'''
```

Sets globally default unit-gain flag value.

```
'''
```

```
cntk_py.set_default_unit_gain_value(value)
```

```
def default_use_mean_gradient_value():
```

[\[docs\]](#)

```
'''
```

Returns true if by default input gradient to learner is averaged.

```
'''
```

```
return cntk_py.default_use_mean_gradient_value()
```

```
def set_default_use_mean_gradient_value(value): [docs]
    '''
    Sets globally default use_mean_gradient_value.
    '''
    cntk_py.set_default_use_mean_gradient_value(value)
```

```
# an internal method to verify that the learning rate schedule
# has a proper (per-sample or per-MB schedule) type and raise
# an exception otherwise
```

```
def _verify_learning_rate_type(learning_rate):
    if not isinstance(learning_rate,
                      (cntk_py.training_parameter_per_sample_schedule,
                       cntk_py.training_parameter_per_minibatch_schedule)):

        raise ValueError('learning_rate type (%s) not supported. '
                          'learning_rate must be a training schedule '
                          '(output of learning_rate_schedule() function)'
                          % type(learning_rate))
```

```
# an internal method to verify that the mometum schedule
# has a proper (per-MB or time-constant schedule) type and raise
# an exception otherwise
```

```
def _verify_momentum_type(momentum):
    if not isinstance(momentum,
                      (cntk_py.training_parameter_per_minibatch_schedule,
                       cntk_py.momentum_as_time_constant_schedule)):

        raise ValueError('momentum type (%s) not supported. '
                          'momentum must be a training schedule '
                          '(output of momentum_schedule() or '
                          'momentum_as_time_constant_schedule() function)'
                          % type(momentum))
```

```
class Learner(cntk_py.Learner): [docs]
    '''
    Abstraction for learning a subset of parameters of a learnable function using first order
    gradient values.
    For example momentum, AdaGrad, RMSProp, etc. are different types of learners with their own
    algorithms for
    learning parameter values using first order gradients.
    To instantiate a concrete learner, use the factory methods in this module.
    '''
```

```
def update(self, gradient_values, training_sample_count): [docs]
    '''
    Update the parameters associated with this learner.

    Args:
        gradient_values (dict): maps :class:`~cntk.variables.Parameter` to
            a NumPy array containing the first order gradient values for the
            Parameter w.r.t. the training objective.
        training_sample_count (int): number of samples in the minibatch

    Returns:
        bool: `False` to indicate that learning has stopped for all of the parameters
        associated with this learner
    '''
    var_nd_map = {var: NDArrayView.from_data(val) for var, val in
                  gradient_values.items()}

    return super(Learner, self)._update(var_nd_map, training_sample_count)
```

```
@property
@typemap
def parameters(self):
    '''
```

```

    The set of parameters associated with this learner.
    """
    return super(Learner, self).parameters()

def reset_learning_rate(self, learning_rate):
    """
    Resets the learning rate. The new schedule is adjusted to be relative
    to the current number of elapsed samples/sweeps: the 0 offset in
    the new schedule corresponds to the current value of elapsed samples/sweeps,
    and it takes effect from the current position in the training process onwards.

    Args:
        learning_rate (output of :func:`learning_rate_schedule`)
        learning_rate to reset to
    """
    _verify_learning_rate_type(learning_rate)
    return super(Learner, self).reset_learning_rate(learning_rate)

def learning_rate(self):
    """
    Current learning rate schedule.
    """
    return super(Learner, self).learning_rate()

class UserLearner(cntk_py.Learner):
    """
    Base class of all user-defined learners. To implement your own learning
    algorithm, derive from this class and override the :meth:`update`.

    Certain optimizers (such as AdaGrad) require additional storage.
    This can be allocated and initialized during construction.
    """

    def __init__(self, parameters, lr_schedule, as_numpy=True):
        super(UserLearner, self).__init__(parameters, lr_schedule)
        self.as_numpy = as_numpy
        self.__disown__()

    def _update(self, gradient_values, training_sample_count, sweep_end):
        """
        Update the parameters and related state associated with this learner.

        Args:
            gradient_values (dict): maps :class:`~cntk.variables.Parameter`
                to a NumPy array containing the gradient for the Parameter w.r.t.
                the training objective.
            training_sample_count (int): number of samples in the minibatch
            sweep_end (bool): if the data is fed by a conforming reader, this
                indicates whether a full pass over the dataset has just occurred.

        Returns:
            bool: `False` to indicate that learning has stopped for all of the
                parameters associated with this learner
        """
        map_if_possible(gradient_values)

        if self.as_numpy:
            var_nd_map = {var: asarray(gradient_values[var]) \
                           for var, val in gradient_values.items()}
        else:
            var_nd_map = gradient_values

        return self.update(gradient_values, training_sample_count, sweep_end)

    def update(self, gradient_values, training_sample_count, sweep_end):
        """
        Update the parameters associated with this learner.

        Args:
            gradient_values (dict): maps :class:`~cntk.variables.Parameter` to
                a NumPy array containing the first order gradient values for the
                Parameter w.r.t. the training objective.
            training_sample_count (int): number of samples in the minibatch

```

```
sweep_end (bool): if the data is fed by a conforming reader, this indicates
whether a full pass over the dataset has just occurred.
```

Returns:

```
bool: `False` to indicate that learning has stopped for all of the
parameters associated with this learner
...
```

```
raise NotImplementedError('UserLearner.update must be overridden')
```

@typemap

```
def training_parameter_schedule(schedule, unit, epoch_size=None):
```

[docs]

```
...
Create a training parameter schedule containing either per-sample (default)
or per-minibatch values.
```

Examples:

```
>>> # Use a fixed value 0.01 for all samples
```

```
>>> s = training_parameter_schedule(0.01, UnitType.sample)
```

```
>>> s[0], s[1]
```

```
(0.01, 0.01)
```

```
>>> # Use 0.01 for the first 1000 samples, then 0.001 for the remaining ones
```

```
>>> s = training_parameter_schedule([0.01, 0.001], UnitType.sample, 1000)
```

```
>>> s[0], s[1], s[1000], s[1001]
```

```
(0.01, 0.01, 0.001, 0.001)
```

```
>>> # Use 0.1 for the first 12 epochs, then 0.01 for the next 15,
```

```
>>> # followed by 0.001 for the remaining ones, with a 100 samples in an epoch
```

```
>>> s = training_parameter_schedule([(12, 0.1), (15, 0.01), (1, 0.001)],
```

```
UnitType.sample, 100)
```

```
>>> s[0], s[1199], s[1200], s[2699], s[2700], s[5000]
```

```
(0.1, 0.1, 0.01, 0.01, 0.001, 0.001)
```

Args:

```
schedule (float or list): if float, is the parameter schedule to be used
for all samples. In case of list, the elements are used as the
values for ``epoch_size`` samples. If list contains pair, the second element is
used as a value for (``epoch_size`` x first element) samples
```

```
unit (:class:`UnitType`): one of two
```

```
* ``sample``: the returned schedule contains per-sample values
```

```
* ``minibatch``: the returned schedule contains per-minibatch values.
```

```
epoch_size (optional, int): number of samples as a scheduling unit.
```

```
Parameters in the schedule change their values every ``epoch_size``
```

```
samples. If no ``epoch_size`` is provided, this parameter is substituted
```

```
by the size of the full data sweep, in which case the scheduling unit is
```

```
the entire data sweep (as indicated by the MinibatchSource) and parameters
```

```
change their values on the sweep-by-sweep basis specified by the
```

```
``schedule``.
```

Returns:

```
training parameter schedule
```

See also:

```
:func:`learning_rate_schedule`
```

```
...
```

```
if unit == UnitType.sample:
```

```
    if isinstance(schedule, cntk_py.training_parameter_per_sample_schedule):
```

```
        return schedule
```

```
else:
```

```
    if isinstance(schedule, cntk_py.training_parameter_per_minibatch_schedule):
```

```
        return schedule
```

```
if isinstance(schedule, (int, float)):
```

```
    if epoch_size is not None:
```

```
        warnings.warn('When providing the schedule as a number, epoch_size is ignored',
```

```
RuntimeWarning)
```

```
    if UnitType(unit) is UnitType.sample:
```

```
        return cntk_py.training_parameter_per_sample_schedule(schedule)
```

```
    else:
```

```
        return cntk_py.training_parameter_per_minibatch_schedule(schedule)
```

```
args = [schedule] if epoch_size is None else [schedule, epoch_size]
```

```
if isinstance(schedule, list):
```

```
    if UnitType(unit) is UnitType.sample:
```

```

        return cntk_py.training_parameter_per_sample_schedule(*args)
    else:
        return cntk_py.training_parameter_per_minibatch_schedule(*args)

    raise ValueError(
        'schedule must be either a float or a list, not %s' % type(schedule))

```

@typemap

```
def learning_rate_schedule(lr, unit, epoch_size=None):
```

[docs]

```

    """
    Create a learning rate schedule (using the same semantics as
    :func:`training_parameter_schedule`).

```

Args:

```

    lr (float or list): see parameter ``schedule`` in
        :func:`training_parameter_schedule`.
    unit (:class:`UnitType`): see parameter
        ``unit`` in :func:`training_parameter_schedule`.
    epoch_size (int): see parameter ``epoch_size`` in
        :func:`training_parameter_schedule`.

```

Returns:

```

    learning rate schedule

```

See also:

```

    :func:`training_parameter_schedule`

```

```

    """
    return training_parameter_schedule(lr, unit, epoch_size)

```

@typemap

```
def momentum_schedule(momentum, epoch_size=None):
```

[docs]

```

    """
    Create a per-minibatch momentum schedule (using the same semantics as
    :func:`training_parameter_schedule` with the `unit=UnitType.minibatch`).

```

Args:

```

    momentum (float or list): see parameter ``schedule`` in
        :func:`training_parameter_schedule`.
    epoch_size (int): see parameter ``epoch_size`` in
        :func:`training_parameter_schedule`.

```

If you want to provide momentum values in a minibatch-size agnostic way, use :func:`momentum_as_time_constant_schedule`.

Examples:

```

>>> # Use a fixed momentum of 0.99 for all samples
>>> m = momentum_schedule(0.99)

>>> # Use the momentum value 0.99 for the first 1000 samples,
>>> # then 0.9 for the remaining ones
>>> m = momentum_schedule([0.99,0.9], 1000)
>>> m[0], m[999], m[1000], m[1001]
(0.99, 0.99, 0.9, 0.9)

>>> # Use the momentum value 0.99 for the first 999 samples,
>>> # then 0.88 for the next 888 samples, and 0.77 for the
>>> # the remaining ones
>>> m = momentum_schedule([(999,0.99),(888,0.88),(0, 0.77)])
>>> m[0], m[998], m[999], m[999+888-1], m[999+888]
(0.99, 0.99, 0.88, 0.88, 0.77)

```

Returns:

```

    momentum schedule

```

```

    """
    return training_parameter_schedule(momentum, UnitType.minibatch, epoch_size)

```

@typemap

```
def momentum_as_time_constant_schedule(momentum, epoch_size=None):
```

[docs]

```

    """
    Create a momentum schedule in a minibatch-size agnostic way
    (using the same semantics as :func:`training_parameter_schedule`

```

```
with `unit=UnitType.sample`).
```

Args:

```
momentum (float or list): see parameter ``schedule`` in
:func:`training_parameter_schedule`.
epoch_size (int): see parameter ``epoch_size`` in
:func:`training_parameter_schedule`.
```

CNTK specifies momentum in a minibatch-size agnostic way as the time constant (in samples) of a unit-gain 1st-order IIR filter. The value specifies the number of samples after which a gradient has an effect of $1/e=37\%$.

If you want to specify the momentum per sample (or per minibatch), use `:func:`momentum_schedule``.

Examples:

```
>>> # Use a fixed momentum of 1100 for all samples
>>> m = momentum_as_time_constant_schedule(1100)

>>> # Use the time constant 1100 for the first 1000 samples,
>>> # then 1500 for the remaining ones
>>> m = momentum_as_time_constant_schedule([1100, 1500], 1000)
```

Returns:

```
momentum as time constant schedule
...
```

```
if isinstance(momentum, (cntk_py.momentum_as_time_constant_schedule)):
    return momentum
```

```
if isinstance(momentum, (int, float)):
    if epoch_size is not None:
        warnings.warn('When providing the schedule as a number, epoch_size is ignored',
```

RuntimeWarning)

```
    return cntk_py.momentum_as_time_constant_schedule(momentum)
```

```
if isinstance(momentum, list):
    args = [momentum] if epoch_size is None else [momentum, epoch_size]
    return cntk_py.momentum_as_time_constant_schedule(*args)
```

```
raise ValueError(
    'momentum must be either a float or a list, not %s' % type(momentum))
```

```
# TODO figure out how to pass inf to C++ in a portable way
```

```
@typemap
```

```
def sgd(parameters, lr,
        l1_regularization_weight=0.0, l2_regularization_weight=0.0,
        gaussian_noise_injection_std_dev=0.0, gradient_clipping_threshold_per_sample=np.inf,
        gradient_clipping_with_truncation=True,
```

[docs]

```
use_mean_gradient=default_use_mean_gradient_value()):
```

```
    '''sgd(parameters, lr, l1_regularization_weight=0, l2_regularization_weight=0,
    gaussian_noise_injection_std_dev=0, gradient_clipping_threshold_per_sample=np.inf,
    gradient_clipping_with_truncation=True)
```

Creates an SGD learner instance to learn the parameters. See [1] for more information on how to set the parameters.

Args:

```
parameters (list of parameters): list of network parameters to tune.
    These can be obtained by the '.parameters()' method of the root
    operator.
lr (output of :func:`learning_rate_schedule`): learning rate schedule.
l1_regularization_weight (float, optional): the L1 regularization weight per sample,
    defaults to 0.0
l2_regularization_weight (float, optional): the L2 regularization weight per sample,
    defaults to 0.0
gaussian_noise_injection_std_dev (float, optional): the standard deviation
    of the Gaussian noise added to parameters post update, defaults to 0.0
gradient_clipping_threshold_per_sample (float, optional): clipping threshold
    per sample, defaults to infinity
gradient_clipping_with_truncation (bool, default ``True``): use gradient clipping
    with truncation
use_mean_gradient (bool, default ``False``): use averaged gradient as input to learner.
    Defaults to the value returned by :func:`default_use_mean_gradient_value`.
```

Returns:

:class:`~cntk.learners.Learner`: learner instance that can be passed to the :class:`~cntk.train.trainer.Trainer`

See also:

[1] L. Bottou. `Stochastic Gradient Descent Tricks`

<<https://www.microsoft.com/en-us/research/publication/stochastic-gradient-tricks>>`_.

Neural

Networks: Tricks of the Trade: Springer, 2012.

...

_verify_learning_rate_type(lr)

gaussian_noise_injection_std_dev = \

training_parameter_schedule(

gaussian_noise_injection_std_dev, UnitType.minibatch)

additional_options = cntk_py.AdditionalLearningOptions()

additional_options.l1_regularization_weight = l1_regularization_weight

additional_options.l2_regularization_weight = l2_regularization_weight

additional_options.gaussian_noise_injection_std_dev = gaussian_noise_injection_std_dev

additional_options.gradient_clipping_threshold_per_sample =

gradient_clipping_threshold_per_sample

additional_options.gradient_clipping_with_truncation = gradient_clipping_with_truncation

additional_options.use_mean_gradient = use_mean_gradient

return cntk_py.sgd_learner(parameters, lr, additional_options)

@typemap

def **momentum_sgd**(parameters, lr, momentum, unit_gain=default_unit_gain_value(),

[docs]

l1_regularization_weight=0.0, l2_regularization_weight=0.0,

gaussian_noise_injection_std_dev=0.0,

gradient_clipping_threshold_per_sample=np.inf,

gradient_clipping_with_truncation=True,

use_mean_gradient=default_use_mean_gradient_value()):

'''momentum_sgd(parameters, lr, momentum, unit_gain=default_unit_gain_value(),

l1_regularization_weight=0.0, l2_regularization_weight=0, gaussian_noise_injection_std_dev=0,

gradient_clipping_threshold_per_sample=np.inf, gradient_clipping_with_truncation=True)

Creates a Momentum SGD learner instance to learn the parameters.

Args:

parameters (list of parameters): list of network parameters to tune.

These can be obtained by the root operator's ``parameters``.

lr (output of :func:`learning_rate_schedule`): learning rate schedule.

momentum (output of :func:`momentum_schedule` or

:func:`momentum_as_time_constant_schedule`): momentum schedule.

For additional information, please refer to the :cntkwiki:`this CNTK Wiki article <BrainScript-SGD-Block#converting-learning-rate-and-momentum-parameters-from-other-toolkits>`.

unit_gain: when ``True``, momentum is interpreted as a unit-gain filter. Defaults to the value returned by :func:`default_unit_gain_value`.

l1_regularization_weight (float, optional): the L1 regularization weight per sample, defaults to 0.0

l2_regularization_weight (float, optional): the L2 regularization weight per sample, defaults to 0.0

gaussian_noise_injection_std_dev (float, optional): the standard deviation of the Gaussian noise added to parameters post update, defaults to 0.0

gradient_clipping_threshold_per_sample (float, optional): clipping threshold per sample, defaults to infinity

gradient_clipping_with_truncation (bool, default ``True``): use gradient clipping with truncation

use_mean_gradient (bool, default ``False``): use averaged gradient as input to learner. Defaults to the value returned by :func:`default_use_mean_gradient_value`.

Returns:

:class:`~cntk.learners.Learner`: learner instance that can be passed to the :class:`~cntk.train.trainer.Trainer`

...

_verify_learning_rate_type(lr)

_verify_momentum_type(momentum)

gaussian_noise_injection_std_dev = \

training_parameter_schedule(

gaussian_noise_injection_std_dev, UnitType.minibatch)

additional_options = cntk_py.AdditionalLearningOptions()

additional_options.l1_regularization_weight = l1_regularization_weight

additional_options.l2_regularization_weight = l2_regularization_weight

additional_options.gaussian_noise_injection_std_dev = gaussian_noise_injection_std_dev


```

        additional_options.gradient_clipping_threshold_per_sample =
gradient_clipping_threshold_per_sample
        additional_options.gradient_clipping_with_truncation = gradient_clipping_with_truncation
        additional_options.use_mean_gradient = use_mean_gradient

    return cntk_py.momentum_sgd_learner(parameters, lr, momentum, unit_gain,
                                        additional_options)

```

@typemap

```

def nesterov(parameters, lr, momentum, unit_gain=default_unit_gain_value(),           [docs]
              l1_regularization_weight=0.0, l2_regularization_weight=0.0,
              gaussian_noise_injection_std_dev=0.0,
              gradient_clipping_threshold_per_sample=np.inf,
              gradient_clipping_with_truncation=True,
              use_mean_gradient=default_use_mean_gradient_value()):
    '''nesterov(parameters, lr, momentum, unit_gain=default_unit_gain_value(),
    l1_regularization_weight=0, l2_regularization_weight=0, gaussian_noise_injection_std_dev=0,
    gradient_clipping_threshold_per_sample=np.inf, gradient_clipping_with_truncation=True)
    Creates a Nesterov SGD learner instance to learn the parameters. This was
    originally proposed by Nesterov [1] in 1983 and then shown to work well in
    a deep learning context by Sutskever, et al. [2].

```

Args:

```

    parameters (list of parameters): list of network parameters to tune.
    These can be obtained by the root operator's ``parameters``.
    lr (output of :func:`learning_rate_schedule`): learning rate schedule.
    momentum (output of :func:`momentum_schedule` or
:func:`momentum_as_time_constant_schedule`): momentum schedule.
    For additional information, please refer to the :cntkwiki:`this CNTK Wiki article
<BrainScript-SGD-Block#converting-learning-rate-and-momentum-parameters-from-other-toolkits>`.
    unit_gain: when ``True``, momentum is interpreted as a unit-gain filter. Defaults
    to the value returned by :func:`default_unit_gain_value`.
    l1_regularization_weight (float, optional): the L1 regularization weight per sample,
    defaults to 0.0
    l2_regularization_weight (float, optional): the L2 regularization weight per sample,
    defaults to 0.0
    gaussian_noise_injection_std_dev (float, optional): the standard deviation
    of the Gaussian noise added to parameters post update, defaults to 0.0
    gradient_clipping_threshold_per_sample (float, optional): clipping threshold
    per sample, defaults to infinity
    gradient_clipping_with_truncation (bool, default ``True``): use gradient clipping
    with truncation
    use_mean_gradient (bool, default ``False``): use averaged gradient as input to learner.
    Defaults to the value returned by :func:`default_use_mean_gradient_value`.

```

Returns:

```

: class:`~cntk.learners.Learner`: learner instance that can be passed to
the : class:`~cntk.train.trainer.Trainer`

```

See also:

[1] Y. Nesterov. A Method of Solving a Convex Programming Problem with Convergence Rate $O(1/\sqrt{k})$. Soviet Mathematics Doklady, 1983.

[2] I. Sutskever, J. Martens, G. Dahl, and G. Hinton. `On the Importance of Initialization and Momentum in Deep Learning` <<http://www.cs.toronto.edu/~fritz/absps/momentum.pdf>>_. Proceedings of the 30th International Conference on Machine Learning, 2013.

...

```

    _verify_learning_rate_type(lr)
    _verify_momentum_type(momentum)
    gaussian_noise_injection_std_dev = \
        training_parameter_schedule(
            gaussian_noise_injection_std_dev, UnitType.minibatch)

    additional_options = cntk_py.AdditionalLearningOptions()
    additional_options.l1_regularization_weight = l1_regularization_weight
    additional_options.l2_regularization_weight = l2_regularization_weight
    additional_options.gaussian_noise_injection_std_dev = gaussian_noise_injection_std_dev
    additional_options.gradient_clipping_threshold_per_sample =
gradient_clipping_threshold_per_sample
    additional_options.gradient_clipping_with_truncation = gradient_clipping_with_truncation
    additional_options.use_mean_gradient = use_mean_gradient

```

```
return cntk_py.nesterov_learner(parameters, lr, momentum, unit_gain,
                                additional_options)
```

@typemap

```
def adadelta(parameters, lr=learning_rate_schedule(1, UnitType.sample), rho=0.95,
epsilon=1e-8,
    l1_regularization_weight=0.0, l2_regularization_weight=0.0,
    gaussian_noise_injection_std_dev=0.0,
    gradient_clipping_threshold_per_sample=np.inf,
    gradient_clipping_with_truncation=True,
    use_mean_gradient=default_use_mean_gradient_value()):
    '''adadelta(parameters, lr, rho, epsilon, l1_regularization_weight=0,
l2_regularization_weight=0, gaussian_noise_injection_std_dev=0,
gradient_clipping_threshold_per_sample=np.inf, gradient_clipping_with_truncation=True)
    Creates an AdaDelta learner instance to learn the parameters. See [1] for
    more information.
```

[docs]

Args:

parameters (list of parameters): list of network parameters to tune.
These can be obtained by the root operator's ``parameters``.
lr (output of :func:`learning_rate_schedule`): learning rate schedule.
rho (float): exponential smooth factor for each minibatch.
epsilon (float): epsilon for sqrt.
l1_regularization_weight (float, optional): the L1 regularization weight per sample,
defaults to 0.0
l2_regularization_weight (float, optional): the L2 regularization weight per sample,
defaults to 0.0
gaussian_noise_injection_std_dev (float, optional): the standard deviation
of the Gaussian noise added to parameters post update, defaults to 0.0
gradient_clipping_threshold_per_sample (float, optional): clipping threshold
per sample, defaults to infinity
gradient_clipping_with_truncation (bool, default ``True``): use gradient clipping
with truncation
use_mean_gradient (bool, default ``False``): use averaged gradient as input to learner.
Defaults to the value returned by :func:`default_use_mean_gradient_value`.

Returns:

:class:`~cntk.learners.Learner`: learner instance that can be passed to
the :class:`~cntk.train.trainer.Trainer`

See also:

[1] Matthew D. Zeiler, `ADADELTA: An Adaptive Learning Rate Method
<<https://arxiv.org/pdf/1212.5701.pdf>>`_.

...

```
gaussian_noise_injection_std_dev = \
    training_parameter_schedule(
        gaussian_noise_injection_std_dev, UnitType.minibatch)
```

```
additional_options = cntk_py.AdditionalLearningOptions()
additional_options.l1_regularization_weight = l1_regularization_weight
additional_options.l2_regularization_weight = l2_regularization_weight
additional_options.gaussian_noise_injection_std_dev = gaussian_noise_injection_std_dev
additional_options.gradient_clipping_threshold_per_sample =
gradient_clipping_threshold_per_sample
additional_options.gradient_clipping_with_truncation = gradient_clipping_with_truncation
additional_options.use_mean_gradient = use_mean_gradient

return cntk_py.ada_delta_learner(parameters, lr, rho, epsilon,
                                additional_options)
```

@typemap

```
def adagrad(parameters, lr, need_ave_multiplier=True,
    l1_regularization_weight=0.0, l2_regularization_weight=0.0,
    gaussian_noise_injection_std_dev=0.0,
    gradient_clipping_threshold_per_sample=np.inf,
    gradient_clipping_with_truncation=True,
    use_mean_gradient=default_use_mean_gradient_value()):
    '''adagrad(parameters, lr, need_ave_multiplier=True, l1_regularization_weight=0,
l2_regularization_weight=0, gaussian_noise_injection_std_dev=0,
gradient_clipping_threshold_per_sample=np.inf, gradient_clipping_with_truncation=True)
    Creates an AdaGrad learner instance to learn the parameters. See [1] for
    more information.
```

[docs]

Args:

parameters (list of parameters): list of network parameters to tune.
 These can be obtained by the root operator's ``parameters``.
 lr (output of :func:`learning_rate_schedule`): learning rate schedule.
 need_ave_multiplier (bool, default):
 l1_regularization_weight (float, optional): the L1 regularization weight per sample, defaults to 0.0
 l2_regularization_weight (float, optional): the L2 regularization weight per sample, defaults to 0.0
 gaussian_noise_injection_std_dev (float, optional): the standard deviation of the Gaussian noise added to parameters post update, defaults to 0.0
 gradient_clipping_threshold_per_sample (float, optional): clipping threshold per sample, defaults to infinity
 gradient_clipping_with_truncation (bool, default ``True``): use gradient clipping with truncation
 use_mean_gradient (bool, default ``False``): use averaged gradient as input to learner. Defaults to the value returned by :func:`default_use_mean_gradient_value`.

Returns:

:class:`~cntk.learners.Learner`: learner instance that can be passed to the :class:`~cntk.train.trainer.Trainer`

See also:

[1] J. Duchi, E. Hazan, and Y. Singer. `Adaptive Subgradient Methods for Online Learning and Stochastic Optimization`
<http://www.magicbroom.info/Papers/DuchiHaSi10.pdf>`. The Journal of Machine Learning Research, 2011.

...

```
_verify_learning_rate_type(lr)
gaussian_noise_injection_std_dev = \
    training_parameter_schedule(
        gaussian_noise_injection_std_dev, UnitType.minibatch)
```

```
additional_options = cntk_py.AdditionalLearningOptions()
additional_options.l1_regularization_weight = l1_regularization_weight
additional_options.l2_regularization_weight = l2_regularization_weight
additional_options.gaussian_noise_injection_std_dev = gaussian_noise_injection_std_dev
additional_options.gradient_clipping_threshold_per_sample =
gradient_clipping_threshold_per_sample
additional_options.gradient_clipping_with_truncation = gradient_clipping_with_truncation
additional_options.use_mean_gradient = use_mean_gradient
```

```
return cntk_py.ada_grad_learner(parameters, lr, need_ave_multiplier,
                                additional_options)
```

@typemap

```
def fsadagrad(parameters, lr, momentum, unit_gain=default_unit_gain_value(),
               variance_momentum=momentum_as_time_constant_schedule(720000),
               l1_regularization_weight=0.0, l2_regularization_weight=0.0,
               gaussian_noise_injection_std_dev=0.0,
               gradient_clipping_threshold_per_sample=np.inf,
               gradient_clipping_with_truncation=True,
               use_mean_gradient=default_use_mean_gradient_value()):
    '''fsadagrad(parameters, lr, momentum, unit_gain=default_unit_gain_value(),
    variance_momentum=momentum_as_time_constant_schedule(720000), l1_regularization_weight=0,
    l2_regularization_weight=0, gaussian_noise_injection_std_dev=0,
    gradient_clipping_threshold_per_sample=np.inf, gradient_clipping_with_truncation=True)
    Creates an FSAdaGrad learner instance to learn the parameters.
```

[docs]

Args:

parameters (list of parameters): list of network parameters to tune.
 These can be obtained by the root operator's ``parameters``.
 lr (output of :func:`learning_rate_schedule`): learning rate schedule.
 momentum (output of :func:`momentum_schedule` or

:func:`momentum_as_time_constant_schedule`): momentum schedule.

For additional information, please refer to the :cntkwiki:`this CNTK Wiki article`
 <BrainScript-SGD-Block#converting-learning-rate-and-momentum-parameters-from-other-toolkits>.

unit_gain: when ``True``, momentum is interpreted as a unit-gain filter. Defaults to the value returned by :func:`default_unit_gain_value`.

variance_momentum (output of :func:`momentum_schedule` or

:func:`momentum_as_time_constant_schedule`): variance momentum schedule. Defaults to ``momentum_as_time_constant_schedule(720000)``.

l1_regularization_weight (float, optional): the L1 regularization weight per sample, defaults to 0.0

l2_regularization_weight (float, optional): the L2 regularization weight per sample, defaults to 0.0

gaussian_noise_injection_std_dev (float, optional): the standard deviation of the Gaussian noise added to parameters post update, defaults to 0.0
 gradient_clipping_threshold_per_sample (float, optional): clipping threshold per sample, defaults to infinity
 gradient_clipping_with_truncation (bool, default ``True``): use gradient clipping with truncation
 use_mean_gradient (bool, default ``False``): use averaged gradient as input to learner. Defaults to the value returned by :func:`default_use_mean_gradient_value`.

Returns:

:class:`~cntk.learners.Learner`: learner instance that can be passed to the :class:`~cntk.train.trainer.Trainer`

...

```

_verify_learning_rate_type(lr)
_verify_momentum_type(momentum)
_verify_momentum_type(variance_momentum)
gaussian_noise_injection_std_dev = \
    training_parameter_schedule(
        gaussian_noise_injection_std_dev, UnitType.minibatch)

```

```

additional_options = cntk_py.AdditionalLearningOptions()
additional_options.l1_regularization_weight = l1_regularization_weight
additional_options.l2_regularization_weight = l2_regularization_weight
additional_options.gaussian_noise_injection_std_dev = gaussian_noise_injection_std_dev
additional_options.gradient_clipping_threshold_per_sample =
gradient_clipping_threshold_per_sample
additional_options.gradient_clipping_with_truncation = gradient_clipping_with_truncation
additional_options.use_mean_gradient = use_mean_gradient

```

```

return cntk_py.fsada_grad_learner(parameters, lr, momentum, unit_gain,
    variance_momentum, additional_options)

```

@typemap

```

def adam(parameters, lr, momentum, unit_gain=default_unit_gain_value(),
    variance_momentum=momentum_as_time_constant_schedule(720000),
    l1_regularization_weight=0.0, l2_regularization_weight=0.0,
    gaussian_noise_injection_std_dev=0.0, gradient_clipping_threshold_per_sample=np.inf,
    gradient_clipping_with_truncation=True,

```

[docs]

```

    use_mean_gradient=default_use_mean_gradient_value(), epsilon=1e-8, adamax=False):

```

```

    '''adam(parameters, lr, momentum, unit_gain=default_unit_gain_value(),
    variance_momentum=momentum_as_time_constant_schedule(720000), l1_regularization_weight=0,
    l2_regularization_weight=0, gaussian_noise_injection_std_dev=0,
    gradient_clipping_threshold_per_sample=np.inf, gradient_clipping_with_truncation=True,
    epsilon=1e-8, adamax=False)

```

Creates an Adam learner instance to learn the parameters. See [1] for more information.

Args:

parameters (list of parameters): list of network parameters to tune. These can be obtained by the root operator's ``parameters``.
 lr (output of :func:`learning_rate_schedule`): learning rate schedule.
 momentum (output of :func:`momentum_schedule` or :func:`momentum_as_time_constant_schedule`): momentum schedule.

For additional information, please refer to the :cntkwiki:`this CNTK Wiki article <BrainScript-SGD-Block#converting-learning-rate-and-momentum-parameters-from-other-toolkits>`.

unit_gain: when ``True``, momentum is interpreted as a unit-gain filter. Defaults to the value returned by :func:`default_unit_gain_value`.

variance_momentum (output of :func:`momentum_schedule` or :func:`momentum_as_time_constant_schedule`): variance momentum schedule. Defaults to ``momentum_as_time_constant_schedule(720000)``.

l1_regularization_weight (float, optional): the L1 regularization weight per sample, defaults to 0.0

l2_regularization_weight (float, optional): the L2 regularization weight per sample, defaults to 0.0

gaussian_noise_injection_std_dev (float, optional): the standard deviation of the Gaussian noise added to parameters post update, defaults to 0.0

gradient_clipping_threshold_per_sample (float, optional): clipping threshold per sample, defaults to infinity

gradient_clipping_with_truncation (bool, default ``True``): use gradient clipping with truncation

use_mean_gradient (bool, default ``False``): use averaged gradient as input to learner. Defaults to the value returned by :func:`default_use_mean_gradient_value`.

epsilon (float, optional): numerical stability constant, defaults to 1e-8

```
adamax: when ``True``, use infinity-norm variance momentum update instead of L2.
Defaults
    to False
```

Returns:

```
:class:`~cntk.learners.Learner`: learner instance that can be passed to
the :class:`~cntk.train.trainer.Trainer`
```

See also:

```
[1] D. Kingma, J. Ba. `Adam: A Method for Stochastic Optimization
<https://arxiv.org/abs/1412.6980>`. International Conference for
Learning Representations, 2015.
```

```
...
```

```
_verify_learning_rate_type(lr)
_verify_momentum_type(momentum)
_verify_momentum_type(variance_momentum)
gaussian_noise_injection_std_dev = \
    training_parameter_schedule(
        gaussian_noise_injection_std_dev, UnitType.minibatch)

additional_options = cntk_py.AdditionalLearningOptions()
additional_options.l1_regularization_weight = l1_regularization_weight
additional_options.l2_regularization_weight = l2_regularization_weight
additional_options.gaussian_noise_injection_std_dev = gaussian_noise_injection_std_dev
additional_options.gradient_clipping_threshold_per_sample =
gradient_clipping_threshold_per_sample
additional_options.gradient_clipping_with_truncation = gradient_clipping_with_truncation
additional_options.use_mean_gradient = use_mean_gradient

return cntk_py.adam_learner(parameters, lr, momentum, unit_gain,
                             variance_momentum, epsilon, adamax, additional_options)
```

```
@typemap
```

```
def rmsprop(parameters, lr,                                     [docs]
            gamma, inc, dec, max, min,
            need_ave_multiplier=True,
            l1_regularization_weight=0.0, l2_regularization_weight=0.0,
            gaussian_noise_injection_std_dev=0.0,
            gradient_clipping_threshold_per_sample=np.inf,
            gradient_clipping_with_truncation=True,
            use_mean_gradient=default_use_mean_gradient_value()):
    '''rmsprop(parameters, lr, gamma, inc, dec, max, min, need_ave_multiplier=True,
    l1_regularization_weight=0, l2_regularization_weight=0, gaussian_noise_injection_std_dev=0,
    gradient_clipping_threshold_per_sample=np.inf, gradient_clipping_with_truncation=True)
    Creates an RMSProp learner instance to learn the parameters.
```

Args:

```
parameters (list of parameters): list of network parameters to tune.
    These can be obtained by the root operator's ``parameters``.
lr (output of :func:`learning_rate_schedule`): learning rate schedule.
gamma (float): Trade-off factor for current and previous gradients. Common value is
0.95. Should be in range (0.0, 1.0)
inc (float): Increasing factor when trying to adjust current learning_rate. Should be
greater than 1
dec (float): Decreasing factor when trying to adjust current learning_rate. Should be
in range (0.0, 1.0)
max (float): Maximum scale allowed for the initial learning_rate. Should be greater
than zero and min
min (float): Minimum scale allowed for the initial learning_rate. Should be greater
than zero
need_ave_multiplier (bool, default ``True``):
l1_regularization_weight (float, optional): the L1 regularization weight per sample,
defaults to 0.0
l2_regularization_weight (float, optional): the L2 regularization weight per sample,
defaults to 0.0
gaussian_noise_injection_std_dev (float, optional): the standard deviation
of the Gaussian noise added to parameters post update, defaults to 0.0
gradient_clipping_threshold_per_sample (float, optional): clipping threshold
per sample, defaults to infinity
gradient_clipping_with_truncation (bool, default ``True``): use gradient clipping
with truncation
use_mean_gradient (bool, default ``False``): use averaged gradient as input to learner.
Defaults to the value returned by :func:`default_use_mean_gradient_value`.
```

Returns:

```

        :class:`~cntk.learners.Learner`: learner instance that can be passed to
        the :class:`~cntk.train.trainer.Trainer`
    """
    _verify_learning_rate_type(lr)
    gaussian_noise_injection_std_dev = \
        training_parameter_schedule(
            gaussian_noise_injection_std_dev, UnitType.minibatch)

    additional_options = cntk_py.AdditionalLearningOptions()
    additional_options.l1_regularization_weight = l1_regularization_weight
    additional_options.l2_regularization_weight = l2_regularization_weight
    additional_options.gaussian_noise_injection_std_dev = gaussian_noise_injection_std_dev
    additional_options.gradient_clipping_threshold_per_sample =
        gradient_clipping_threshold_per_sample
    additional_options.gradient_clipping_with_truncation = gradient_clipping_with_truncation
    additional_options.use_mean_gradient = use_mean_gradient

    return cntk_py.rmsprop_learner(parameters, lr, gamma, inc, dec, max, min,
                                    need_ave_multiplier, additional_options)

```

@typemap

```
def universal(update_func, parameters):
```

[docs]

Creates a learner which uses a CNTK function to update the parameters.

Args:

update_func: function that takes parameters and gradients as arguments and returns a :class:`~cntk.ops.functions.Function` that performs the desired updates. The returned function updates the parameters by means of containing :func:`~cntk.ops.assign` operations. If ``update_func`` does not contain :func:`~cntk.ops.assign` operations the parameters will not be updated.
 parameters (list): list of network parameters to tune. These can be obtained by the root operator's `parameters`.

Returns:

:class:`~cntk.learners.Learner`: learner instance that can be passed to the :class:`~cntk.train.trainer.Trainer`

Examples:

```

>>> def my_adagrad(parameters, gradients):
...     accumulators = [C.constant(0, shape=p.shape, dtype=p.dtype, name='accum') for p
in parameters]
...     update_funcs = []
...     for p, g, a in zip(parameters, gradients, accumulators):
...         accum_new = C.assign(a, g * g)
...         update_funcs.append(C.assign(p, p - 0.01 * g / C.sqrt(accum_new + 1e-6)))
...     return C.combine(update_funcs)
...
>>> x = C.input_variable((10,))
>>> y = C.input_variable((2,))
>>> z = C.layers.Sequential([C.layers.Dense(100, activation=C.relu),
C.layers.Dense(2)])(x)
>>> loss = C.cross_entropy_with_softmax(z, y)
>>> learner = C.universal(my_adagrad, z.parameters)
>>> trainer = C.Trainer(z, loss, learner)
>>> # now trainer can be used as any other Trainer
...

```

```

from .. import constant
args, _ = utils.get_python_function_arguments(update_func)
if len(args) != 2:
    raise ValueError('update_func must be a function that accepts two arguments
(parameters, gradients)')
gradients = []
for p in parameters:
    if any(dim<0 for dim in p.shape):
        raise ValueError('parameter %s has inferred dimensions. Please create the learner
after all parameter shapes have been determined'%str(p))
    gradients.append(constant(0, shape=p.shape, dtype=p.dtype, name='grad'))

result = update_func(parameters, gradients)

return cntk_py.universal_learner(parameters, gradients, result)

```

