

Stack Overflow is a community of 4.7 million programmers, just like you, helping each other.

Join them; it only takes a minute:

Sign up

Join the Stack Overflow community to:


Ask programming questions

Answer and help your peers

Get recognized for your expertise

Python (NumPy, SciPy), finding the null space of a matrix

Work on work you love. From home.



I'm trying to find the null space (solution space of $Ax=0$) of a given matrix. I've found two examples, but I can't seem to get either to work. Moreover, I can't understand what they're doing to get there, so I can't debug. I'm hoping someone might be able to walk me through this.

The documentation pages ([numpy.linalg.svd](#) , and [numpy.compress](#)) are opaque to me. I learned to do this by creating the matrix $C = [A|0]$, finding the reduced row echelon form and solving for variables by row. I can't seem to follow how it's being done in these examples.

Thanks for any and all help!

Here is my sample matrix, which is the same as the [wikipedia example](#):

```
A = matrix([
    [2,3,5],
    [-4,2,3]
])
```

Method ([found here](#), and [here](#)):

```
import scipy
from scipy import linalg, matrixr
def null(A, eps=1e-15):
    u, s, vh = scipy.linalg.svd(A)
    null_mask = (s <= eps)
    null_space = scipy.compress(null_mask, vh, axis=0)
    return scipy.transpose(null_space)
```

When I try it, I get back an empty matrix:

```
Python 2.6.6 (r266:84292, Sep 15 2010, 16:22:56)
[GCC 4.4.5] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> import scipy
>>> from scipy import linalg, matrix
>>> def null(A, eps=1e-15):
...     u, s, vh = scipy.linalg.svd(A)
...     null_mask = (s <= eps)
...     null_space = scipy.compress(null_mask, vh, axis=0)
...     return scipy.transpose(null_space)
...
>>> A = matrix([
...     [2,3,5],
...     [-4,2,3]
... ])
>>>
>>> null(A)
array([], shape=(3, 0), dtype=float64)
>>>
```

[numpy](#) [matrix](#) [scipy](#) [linear-algebra](#) [svd](#)

edited Nov 9 '15 at 22:21



[ali_m](#)

24.4k

4

51

100

asked May 4 '11 at 19:56



[Nona Urbiz](#)

1,724

11

38

73

-
- 2 The wikipedia page you linked to actually gives a very nice explanation of why you should use an SVD to calculate the null space (or solve) of a matrix when you're dealing with floating point values. [en.wikipedia.org/wiki/...](http://en.wikipedia.org/wiki/Null_space) The approach you describe (solving for variables row-by-row) will amplify any rounding errors, etc. (This is the same reason you should almost never explicitly calculate the inverse of a matrix...) – [Joe Kington](#) May 4 '11 at 20:29
-

5 Answers

It appears to be working okay for me:

```
A = matrix([[2,3,5],[-4,2,3],[0,0,0]])
A * null(A)
>>> [[ 4.02455846e-16]
>>> [ 1.94289029e-16]
>>> [ 0.00000000e+00]]
```

answered May 4 '11 at 20:29



[Bashwork](#)

1,366 6 13

I'm sure I'm missing something, but Wikipedia suggests that the values should be $\begin{bmatrix} -0.0625 \\ -1.625 \\ 1 \end{bmatrix}$? – [Nona Urbiz](#) May 4 '11 at 20:36

Moreover, it's returning an empty matrix for me $\begin{bmatrix} \end{bmatrix}$. What could be wrong? – [Nona Urbiz](#) May 4 '11 at 20:58

6 @Nona Urbiz - It's returning an empty matrix because you're not putting in a row of zeros, as Bashwork (and wikipedia) does above. Also, the null space values returned ($\begin{bmatrix} -0.33 \\ -0.85 \\ 0.52 \end{bmatrix}$) are normalized so that the magnitude of the vector is 1. The wikipedia example is not normalized. If you just take $n = \text{null}(A)$ and have a look at $n / n.\text{max}()$, you'll get $\begin{bmatrix} -0.0625 \\ -1.625 \\ 1 \end{bmatrix}$. – [Joe Kington](#) May 4 '11 at 21:08

3 @Bashwork - How would I know to programmatically add a row of zeroes? Does the matrix have to be square? – [Coder](#) Sep 2 '12 at 20:27



Did you find this question interesting? Try our newsletter

Sign up for our newsletter and get our top new questions delivered to your inbox ([see an example](#)).

You get the SVD decomposition of the matrix A . s is a vector of eigenvalues. You are interested in almost zero eigenvalues (see $Ax = \lambda x$ where $|\lambda| < \epsilon$), which is given by the vector of logical values `null_mask`.

Then, you extract from the list `vh` the eigenvectors corresponding to the almost zero eigenvalues, which is exactly what you are looking for: a way to span the null space. Basically, you extract the rows and then transpose the results so that you get a matrix with eigenvectors as columns.

edited May 4 '11 at 20:17

answered May 4 '11 at 20:05



Wok

2,786 2 17 41

Thank you very much for taking the time to reply and help me. Your answer was very helpful to me, but I accepted Bashworks answer as ultimately, it brought me to the solution. The only reason I am able to understand the solution though, is your response. – [Nona Urbiz](#) May 4 '11 at 21:16

No worry, I thought your problem was something else. – [Wok](#) May 5 '11 at 10:03

Sympy makes this straightforward.

```
>>> from sympy import Matrix
>>> A = [[2, 3, 5], [-4, 2, 3], [0, 0, 0]]
>>> A = Matrix(A)
>>> A * A.nullspace()[0]
Matrix([
[0],
[0],
[0]])
>>> A.nullspace()
[Matrix([
[-1/16],
[-13/8],
[ 1]])]
```

edited Nov 11 '15 at 2:16

answered Nov 10 '14 at 2:09



ldr

1,447 2 16 34

Your method is almost correct. The issue is that the shape of `s` returned by the function `scipy.linalg.svd` is $(K,)$ where $K=\min(M,N)$. Thus, in your example, `s` only has two entries (the

singular values of the first two singular vectors). The following correction to your null function should allow it to work for any sized matrix.

```
import scipy
import numpy as np
from scipy import linalg, matrix
def null(A, eps=1e-12):
    ... u, s, vh = scipy.linalg.svd(A)
    ... padding = max(0, np.shape(A)[1]-np.shape(s)[0])
    ... null_mask = np.concatenate(((s <= eps), np.ones((padding,), dtype=bool)), axis=0)
    ... null_space = scipy.compress(null_mask, vh, axis=0)
    ... return scipy.transpose(null_space)
A = matrix([[2,3,5],[-4,2,3]])
print A*null(A)
>>>[[ 4.44089210e-16]
>>> [ 6.66133815e-16]]
A = matrix([[1,0,1,0],[0,1,0,0],[0,0,0,0],[0,0,0,0]])
print null(A)
>>>[[ 0.         -0.70710678]
>>> [ 0.          0.         ]
>>> [ 0.         0.70710678]
>>> [ 1.          0.         ]]
print A*null(A)
>>>[[ 0.  0.]
>>> [ 0.  0.]
>>> [ 0.  0.]
>>> [ 0.  0.]
```

edited Dec 19 '14 at 16:56

answered Oct 23 '14 at 19:25



Thomas Wentworth

31 4

I have been using this code in my work and I noticed a problem. An eps value of 1e-15 seems to be too small. Notably, consider the matrix $A = \text{np.ones}(13,2)$. This code will report that this matrix has a rank 0 null space. This is due to the `scipy.linalg.svd` function reporting that the second singular value is above 1e-15. I don't know much about the algorithms behind this function, however I suggest using `eps=1e-12` (and perhaps lower for very large matrices) unless someone with more knowledge can chime in. (In infinite precision the second singular value should be 0). – Thomas Wentworth Dec 19 '14 at 16:55

A faster but less numerically stable method is to use a [rank-revealing QR decomposition](#), such as `scipy.linalg.qr` with `pivoting=True` :

```
import numpy as np
from scipy.linalg import qr

def qr_null(A, tol=None):
    Q, R, P = qr(A.T, mode='full', pivoting=True)
    tol = np.finfo(R.dtype).eps if tol is None else tol
    rnk = min(A.shape) - np.abs(np.diag(R))[::-1].searchsorted(tol)
    return Q[:, rnk:].conj()
```

For example:

```
A = np.array([[ 2, 3, 5],
              [-4, 2, 3],
              [ 0, 0, 0]])
Z = qr_null(A)

print(A.dot(Z))
#[[ 4.44089210e-16]
# [ 6.66133815e-16]
# [ 0.00000000e+00]]
```

edited Nov 10 '15 at 2:19

answered Nov 9 '15 at 22:15



[ali_m](#)

24.4k

4

51

100