Discussion

sandipan_dey *Student*

17 discussions started 89 comments

Show all posts ▼ by recent activity ▼

anyone interested in future collaboration? 10

Hi, I enjoyed this course and the discussions have been helpful for me time to time. I am interested in pursuing further by working on some real ...

Pinned

[Note from Staff] 1.2 has been regraded 22

Indeed Problem 1.2 contained a grader error for the plot depicting the linear classifier. It has now been regraded. Please let us know if there are...

Staff

[Staff] Answers to 1. (2) look incorrect 28

[STAFF] Confusion regarding Q 2 (2)/(3) 1

Following

? [Staff] Please tell me my total grade
Could you please tell me my total grade as I didn't do the exam yet, so to know If I now passed the course. I'm so nervous. Thank you
11

Optimal episodic reward (thought process)
48

? Struggling with Linear Q-learning
Where an I going wrong? q_value = (theta @ current_state_vector)[tuple2index(action_index, object_index)] max_q_value_next = 0 if terminal else...
30

? [Staff] Evaluate linear Q-learning on Home World game
Community TA
26

[Staff] Report performance
43

Q1: Does highest possible salary here refer to highest expected salary?
Q1: Just wanted to confirm whether "highest possible salary" here refers to highest expected salary.(we are supposed to consider the risk or T al...
Following
12

? [Staff] Epsilon-greedy method 2
Is question referring to epsilon at the beginning of the learning or end of the learning or is it a fixed value?
11

All Posts

[STAFF] Confusion regarding Q 2 (2)/(3)

discussion posted about 9 hours ago by **sandipan_dey**

Just reading the '*ideal policy*' I just assumed (wrongly) that the question is asking for the *ideal policy* at *convergence*, since

1. we already computed the q-value matrix for the *first iteration* in part (1), i thought now we need to go for the values at *convergence* and missed the part that it was asked for the *first iteration* again.
2. also the values *first iteration* are already computed in (1), they can trivially be obtained from the computed matrix in (1), i thought they won't be asked again for part (3).

3. given the difficulty level of the exam I never thought that same question could be asked twice.

The *ideal policy* at *convergence* remains same as that after the *first iteration* but the $V_1\ (s = 0)$ and $V_1\ (s = 1)$ values are different *at convergence* (3.25 and 6.5, instead of 0 and 4, respectively),.

I know it's my bad (did not read the question properly), but given that I did some over-work (and obtained the final values at *convergence* correctly) could I be considered for some partial marks :-)?

This post is visible to everyone.

0 responses

Preview

Submit