


EdX and its Members use cookies and other tracking technologies for performance, analytics, and marketing purposes. By using this website, you accept this use. Learn more about these technologies in the [Privacy Policy](#).



[Course](#) > [Midterm Exam \(1 week\)](#) > [Midterm Exam 1](#) > Problem 5

## Problem 5


Consider a 2-layer feed-forward neural network that takes in  $x \in \mathbb{R}^2$  and has two ReLU hidden units as defined in the figure below. **Note that hidden units have no offset parameters in this problem.**



## 5. (1)

4/4 points (graded)

The values of the weights in the hidden layer are set such that they result in the  $z_1$  and  $z_2$  "classifiers" as shown in the  $(x_1, x_2)$ -space in the figure below:



The  $z_1$  "classifier" with the normal  $w_1 = [w_{11} \ w_{21}]^T$  is the line given by  $z_1 = x \cdot w_1 = 0$ .  
 Similarly, the  $z_2$  "classifier" with the normal  $w_2 = [w_{12} \ w_{22}]^T$  is the line given by  $z_2 = x \cdot w_2 = 0$ .

The arrows labeled  $w_1$  and  $w_2$  point in the **positive** directions of the respective normal vectors.

**The regions labeled I, II, III, IV are the 4 regions defined by these two lines not including the boundaries.**

Choose the region(s) in  $(x_1, x_2)$  space which are mapped into each of the following regions in  $(f_1, f_2)$ -space, the 2-dimensional space of hidden unit activations  $f(z_1)$  and  $f(z_2)$ . (For example, for the second column below, choose the region(s) in  $(x_1, x_2)$  space which are mapped into the  $f_1$ -axis in  $(f_1, f_2)$ -space.)

(Choose all that apply for each column.)

$\{(f_1, f_2) : f_1 > 0, f_2 > 0\}$  :  $f_1$ -axis:

$f_2$ -axis:

the origin  $(f_1, f_2) = (0, 0)$ :

(Choose all that apply.)

<input checked="" type="checkbox"/> I ✓	<input type="checkbox"/> I	<input type="checkbox"/> I	<input type="checkbox"/> I
<input type="checkbox"/> II	<input checked="" type="checkbox"/> II ✓	<input type="checkbox"/> II	<input type="checkbox"/> II
<input type="checkbox"/> III	<input checked="" type="checkbox"/> III ✓	<input checked="" type="checkbox"/> III ✓	<input checked="" type="checkbox"/> III ✓
<input type="checkbox"/> IV	<input type="checkbox"/> IV	<input checked="" type="checkbox"/> IV ✓	<input type="checkbox"/> IV
<input type="checkbox"/> None of the above	<input type="checkbox"/> None of the above	<input type="checkbox"/> None of the above	<input type="checkbox"/> None of the above
✓	✗	✗	✓


*Correction Note (July 30 03:00UTC):* In an earlier version, the problem statement did not include the emphasis "in the  $(x_1, x_2)$ -space " in the first sentence.

*Correction Note (August 1 13:00UTC):* In an earlier version, the caption under the figure did not include "not including the boundaries".

*Correction Note (August 4 04:00UTC):* Added "(For example, for the second column below, choose the region(s) in  $(x_1, x_2)$  space which are mapped into the  $f_1$ -axis in  $(f_1, f_2)$ -space.)"

**Grading Note:** Since we have separately asked for the region(s) that map(s) into the origin (in the fourth column), and we do not want to penalize on the interpretation on with axes include the origin, we will accept the region **II** as the answer for the  $f_1$ -axis and region **IV** as the answer for the  $f_2$ -axis.

**Solution:**



The regions I, II, III, IV are defined by (but do not include) the lines  $z_1 = x \cdot w_1 = 0$  and  $z_2 = x \cdot w_2 = 0$ . Hence,

$$z_1 \begin{cases} > 0 & \text{in I, II} \\ < 0 & \text{in I, II} \end{cases}$$

$$z_2 \begin{cases} > 0 & \text{in I, IV} \\ < 0 & \text{in II, III} \end{cases}$$

Applying the reLu function, we get

$$f_1 = f(z_1) \begin{cases} > 0 & \text{in I, II} \\ = 0 & \text{in III, IV} \end{cases}$$

$$f_2 = f(z_2) \begin{cases} > 0 & \text{in I, IV} \\ = 0 & \text{in II, III} \end{cases}.$$

Hence

- The region I in  $(x_1, x_2)$ -space maps into the region  $\{(f_1, f_2) : f_1 > 0, f_2 > 0\}$  : in  $(f_1, f_2)$ -space ;
- The regions II,III maps into the region  $\{(f_1, f_2) : f_1 \geq 0, f_2 = 0\}$ , which is the  $f_1$ -axis in  $(f_1, f_2)$ -space ;
- The regions III,IV maps into the region  $\{(f_1, f_2) : f_2 \geq 0, f_1 = 0\}$ , which is the  $f_2$ -axis in  $(f_1, f_2)$ -space ;
- The regions III maps to  $\{(f_1, f_2) : f_2 = 0, f_1 = 0\}$ , the origin in  $(f_1, f_2)$ -space.

Submit

You have used 2 of 3 attempts

**i** Answers are displayed within the problem

## 5. (2)

1/2 points (graded)

If we keep the hidden layer parameters above fixed but add and train additional hidden layers (applied after this layer) to further transform the data, could the resulting neural network solve this classification problem?

☒ yes ✖☐ no ✔

Suppose we stick to the 2-layer architecture but add many more ReLU hidden units, all of them without offset parameters. Would it be possible to train such a model to perfectly separate these points?

**Grading Note:** Since there is ambiguity as to whether some data points lie on the same line through the origin, we will accept both answers as correct.

☐ yes ✔☒ no ✔

**Solution:**

- Since points with different labels, namely those in region **III**, are all mapped to the origin, it is impossible to classify these correctly by only adding more hidden layers.
- The points in regions **I,II, IV** can be assumed to be distinguished by either the  $f_1$  or  $f_2$  coordinates, so we only need to separate the points in region **III**, and adding more units, which correspond to more lines through the origin, will work.

Submit

You have used 1 of 3 attempts

**i** Answers are displayed within the problem

## 5. (3)

5/5 points (graded)

Which of the following statements is correct?

1. The gradient calculated in the backpropagation algorithm consists of the partial derivatives of the loss function with respect to each network weight.

☒ True ✓☐ False

2. Initialization of the parameters is often important when training large feed-forward neural networks.

If weights in a neural network with sigmoid units are initialized all the weights to close to zero values, then during early stochastic gradient descent steps, the network represents a nearly linear function of the inputs.

☒ True ✓

☐ False

3. On the other hand, if we randomly set all the weights to very large values, or don't scale them properly with the number of units in the layer below, then the sigmoid units would behave like sign units.

(Note that a sign unit is a unit with activation function  $\text{sign}(x) = 1$  if  $x > 0$  and  $\text{sign}(x) = -1$  if  $x < 0$ . For the purpose of this question, it does not matter what  $\text{sign}(0)$  is.)

☐ True ✓

☒ False ✓

**Grading Note:** Since the question did not specify whether "behave like sign units" allows for shifting or rescaling of the sign function, both "True" and "False" are accepted as correct.

4. If we use only sign units in a feedforward neural network, then the stochastic gradient descent update will



☒ almost never change any of the weights ✓

☐ change the weights by large amounts at random

5. Stochastic gradient descent differs from (true) gradient descent by updating only one network weight during each gradient descent step.

☐ True

☒ False ✓

*Correction note (July 31 16:00UTC):*. In the earlier version, the sign unit definition was not included.

**Solution:**

- True, by definition of the backpropagation algorithm.
- True, because those activation functions are linear-like near zero.
- True, because far from zero the sigmoid function looks like a scaled and shifted version of the sign function.
- True, because of the gradient is zero except at one point.

- False. Stochastic gradient descent differs by considering the gradient with respect to just one training example on each step.

You have used 2 of 3 attempts

**i** Answers are displayed within the problem

## 5. (4)

3/3 points (graded)

There are many good reasons to use convolutional layers in CNNs as opposed to replacing them with fully connected layers. Please check T or F for each statement.

Since we apply the same convolutional filter throughout the image, we can learn to recognize the same feature wherever it appears.

☒ True ✓

☐ False

A fully connected layer for a reasonably sized image would simply have too many parameters

☒ True ✓

☐ False

A fully connected layer can learn to recognize features anywhere in the image even if the features appeared preferentially in one location during training

☐ True

☒ False ✓

### Solution:

- True by definition
- True. For a fully connected layer of size  $n$  and a picture of size  $w \times h$ , the size of the layer would be  $whn$ , several times larger than the image itself.
- False. Only a convolutional layer can do this. A fully connected layer's weights depend on the location, so they would only be able to recognize features at locations that appeared preferentially in the training data.

Submit

You have used 1 of 3 attempts

**i** Answers are displayed within the problem

## Error and Bug Reports/Technical Issues

[Hide Discussion](#)






**Topic:** Midterm Exam (1 week):Midterm Exam 1 / Problem 5

[Add a Post](#)

Show all posts ▼

by recent activity ▼

- |   |   |    |
|---|---|----|
| ? | <a href="#">[STAFF] 5.(2) second question: Is grader's answer incorrect, because there are points with different labels lying on the same line through the origin?</a>  | 7  |
| ✓ | <a href="#">[Staff] Problem 5. (1): Do f1-axis and f2-axis include the origin or not?</a>   | 12 |
| 💬 | <a href="#">[Staff] Question 5.(3).1 Answer</a><br>The original statement: "The gradient calculated in the backpropagation algorithm consists of the partial derivatives of the loss function with res..."                        | 9  |
| 💬 | <a href="#">[STAFF] The image at q 5.1 is another now, is it OK?</a><br>Positions of the points at the image are slightly different with the one I seen when I take the exam. I'm sure about this as I sent question to staff ... | 18 |
| ✓ | <a href="#">[TA&amp;Staff] 5.2 solution and explanation.</a>  | 7  |
| ? | <a href="#">[STAFF] 5.3.4 gradient of sign isn't 0 it's <math>2 * \text{dirac delta}</math> for <math>x \sim 0</math> which is very large for values close to 0</a><br>Can I get clarification on this question?                  | 1  |
| ? | <a href="#">[Staff] Solution of 5.3.1</a><br>In my understanding the gradient calculated in the backpropagation algorithm besides of the partial derivatives of the loss function with respect ...                                | 1  |
| 💬 | <a href="#">[Midterm] Nice exam</a><br>I planned to this exam with time and used it to study as if it was a homework. I was smashed by Problem 5, but some of the questions I thought r...  | 2  |

 <u>Midterm Question 5.1</u>	17
<u>I was careless as I omitted the origin from the f1 and f2 axes....</u>	
 <u>Ugh, (0,0) is also on axes</u>	2
<u>shouldn't have lost these two points. Maybe we should get partial credits?</u>	
 <u>question on the problem 5.2 second part</u>	1
<u>In the region II, two lower points with DIFFERENT label sit on the same line that passes through the ORIGIN. How you can distinguish these two p...</u>	
 <u>[Staff] 5.1 typo in solution</u>	1
 <u>[STAFF] Need clarification on question 5.(2) 2 solution</u>	1
<u>The solution is: "The points in regions I, II, IV can be assumed to be distinguished by either the f1 or f2 coordinates." Is it enough if we have 2 la...</u>	

© All Rights Reserved