

EdX and its Members use cookies and other tracking technologies for performance, analytics, and marketing purposes. By using this website, you accept this use. Learn more about these technologies in the [Privacy Policy](#).



MITx: 6.86x

Machine Learning with Python-From Linear Models to Deep Learning

[Help](#)[sandipan\\_dey](#)

[Unit 2 Nonlinear Classification](#),  
[Linear regression, Collaborative](#)

[Course](#) > [Filtering \(2 weeks\)](#)

> [Project 2: Digit recognition \(Part 1\)](#) > 3. Support Vector Machine

### 3. Support Vector Machine

Bob thinks it is clearly not a regression problem, but a classification problem. He thinks that we can change it into a binary classification and use the support vector machine we learned in Lecture 4 to solve the problem. In order to do so, he suggests that we can build an one vs. rest model for every digits. For example, classifying the digits into two classes: 0 and not 0.

Bob wrote a function `run_svm_one_vs_rest_on_MNIST` where he changed the labels of digit 1-9 to 1 and keeps the label 0 for digit 0. He also found that `sklearn` package contains a SVM model where you can use directly. He gave you the link to that model and hope you can tell him how to use that.

**You will be working in the file `part1/svm.py` in this problem**

**Important:** For this problem, you will need to use the [scikit-learn](#) library. If you don't have it, install it using `pip install sklearn`

#### One vs. Rest SVM

5/5 points (graded)

Use the `sklearn` package and build the SVM model on your local machine. Use `random_state = 0`, `C=0.1` and default values for other parameters.

**Available Functions:** You have access to the `sklearn`'s implementation of the linear SVM as `LinearSVC`; No need to import anything.

```
1 def one_vs_rest_svm(train_x, train_y, test_x):
2     """
3     Trains a linear SVM for binary classification
4
5     Args:
6     train_x - (n, d) NumPy array (n datapoints each with d features)
7     train_y - (n, ) NumPy array containing the labels (0 or 1) for each training data point
8     test_x - (m, d) NumPy array (m datapoints each with d features)
```

```
9 Returns:
10     pred_test_y - (m,) NumPy array containing the labels (0 or 1) for each test data point
11     """
12     clf = LinearSVC(random_state = 0, C=0.1)
13     clf.fit(train_x, train_y)
14     return clf.predict(test_x)
15
```

Press ESC then TAB or click outside of the code editor to exit

Correct

## Test results

**CORRECT**

Test: output1

Output:

Hide output

```
train_x: [[0.74021551 0.43195563]
 [0.40684059 0.44127392]
 [0.99915919 0.52231402]
 [0.29457754 0.45500002]
 [0.25464916 0.50020259]
 [0.54634547 0.96572388]
 [0.27012457 0.04665474]
 [0.76504382 0.82860708]
 [0.68958708 0.77897273]
 [0.38031509 0.81979835]
 [0.84065229 0.79751435]
 [0.95114222 0.29275591]
 [0.06657971 0.70887626]
 [0.37758739 0.57958282]
 [0.59278639 0.24966864]
 [0.46846513 0.82311161]]
train_y: [0 0 0 0 0 0 0 1 1 1 0 1 1 1 0 0]
test_x: [[0.08559828 0.35649492]
 [0.58029473 0.48915356]
 [0.95243623 0.31967552]
 [0.3185904 0.50873137]
 [0.48851915 0.9699501 ]
 [0.2590593 0.80253755]
 [0.73428524 0.88636573]
 [0.54701302 0.1022148 ]
 [0.10810793 0.8151251 ]
 [0.97802428 0.58690018]
 [0.46722661 0.55904237]]
Submission output: [0 0 0 0 0 0 0 0 0 0 0]
```

Test: output2

**Output:**

```
train_x: [[0.65447353 0.20028434 0.48239213]
[0.64171897 0.56762668 0.88175238]
[0.8447331 0.4998508 0.90182049]
[0.63844322 0.55868315 0.5755569 ]
[0.24214408 0.51659304 0.76089543]
[0.54918455 0.26628602 0.22222436]
[0.2604621 0.47727055 0.97995404]
[0.63602578 0.42497548 0.32957748]
[0.72721527 0.3376802 0.32333993]
[0.5883944 0.72746542 0.62569041]
[0.72004228 0.52763629 0.71935559]]
train_y: [0 0 1 0 1 0 1 0 0 1 1]
test_x: [[0.77964731 0.81920037 0.25430089]
[0.66557967 0.78041885 0.28126548]
[0.0423602 0.02576799 0.57016067]
[0.13564713 0.7721141 0.34966991]
[0.15671801 0.3125784 0.30351681]
[0.33528671 0.38876823 0.19711744]
[0.28158124 0.20749997 0.24147398]
[0.2879429 0.07601778 0.26567582]
[0.28807974 0.57639994 0.42400284]
[0.53160431 0.03197762 0.89518054]
[0.23845549 0.73092687 0.44063511]
[0.74996488 0.64791171 0.08330063]
[0.30355642 0.72450119 0.75979937]
[0.55121859 0.60209421 0.06778405]
[0.65039024 0.35902444 0.78136871]]
Submission output: [0 0 1 1 0 0 0 0 0 1 1 0 1 0 0]
```

Test: output3

Output:

```
train_x: [[8.11279050e-01 4.52141831e-01]
 [2.10601269e-01 4.60914351e-01]
 [7.56697823e-01 1.80289453e-01]
 [7.06835399e-01 3.74637157e-01]
 [8.38588837e-01 6.44637236e-01]
 [4.74836638e-01 7.96392933e-04]
 [8.62446759e-01 8.94038847e-01]
 [3.79948573e-01 3.09683287e-02]
 [6.67359468e-03 9.36768967e-01]
 [8.78075213e-01 1.76902314e-01]
 [4.77391055e-01 4.33192373e-01]
 [4.43228794e-01 1.01555361e-02]
 [5.76255389e-01 7.91439018e-01]
 [9.68005407e-01 5.91789850e-01]
 [5.14374622e-01 4.40404097e-01]
 [6.03612168e-01 6.45366923e-02]
 [2.19506810e-01 9.68929440e-01]
 [8.73675061e-01 7.23780916e-02]
 [2.21873555e-01 8.56756712e-02]]
train_y: [0 0 0 1 1 1 1 1 1 0 1 0 0 0 0 1 1 1 0]
test_x: [[0.81382256 0.15376929]
 [0.23824971 0.33176453]
 [0.02746353 0.83300153]
 [0.17765254 0.40144322]
 [0.32422425 0.18189357]
 [0.87095893 0.57074161]
 [0.67063566 0.14012018]
 [0.11496733 0.14235777]
 [0.96174966 0.90831115]
 [0.83568426 0.16983058]
 [0.42290956 0.21958185]
 [0.12218531 0.46372719]
 [0.60522441 0.29051125]
 [0.07373684 0.62676753]
 [0.20730928 0.7869615 ]
 [0.35580997 0.46428361]
 [0.78131668 0.69473551]
 [0.87077188 0.91518852]]
Submission output: [0 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1]
```

Test: outputr4

Output:

```
train_x: [[0.40818312 0.38159886 0.24241605 0.38366628]
 [0.27033749 0.00799247 0.68875155 0.90568713]
 [0.7260674 0.45649055 0.18981836 0.21927554]
 [0.86720485 0.37834952 0.12136992 0.04190428]
 [0.76908809 0.13222895 0.75322116 0.57165623]
 [0.2783364 0.08799754 0.44468719 0.14840139]
 [0.6340726 0.84103611 0.87859193 0.58786321]
 [0.75026416 0.28905957 0.85672734 0.26242404]
 [0.51058753 0.874677 0.64296309 0.74521429]
 [0.21944755 0.75669826 0.65450055 0.07229256]
 [0.00590142 0.57580201 0.04343592 0.64508964]
 [0.55628962 0.60103265 0.32649291 0.44102935]]
train_y: [1 1 1 0 1 0 0 1 1 1 0 0]
test_x: [[0.37860324 0.13043979 0.81896389 0.26955548]
 [0.71999832 0.0391397 0.0892016 0.00906736]
 [0.80751849 0.84190127 0.63861565 0.20653685]
 [0.44203692 0.51044744 0.11300264 0.96149892]
 [0.55899333 0.30114035 0.65591599 0.02560668]
 [0.43083021 0.6259706 0.56773372 0.2968281 ]
 [0.16241506 0.59334928 0.35972919 0.55718829]
 [0.93006913 0.88344789 0.23734165 0.15495788]
 [0.67133825 0.23542507 0.25029202 0.22982249]
 [0.6335627 0.66607768 0.62857216 0.64588296]
 [0.67391828 0.13442488 0.51791969 0.62855356]
 [0.83447665 0.33478579 0.7688467 0.82935934]
 [0.2353971 0.62110565 0.55023205 0.09117792]
 [0.61131201 0.80061459 0.54041534 0.86326114]
 [0.0276005 0.62259065 0.59953523 0.31203471]]
Submission output: [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1]
```

Test: size

Output:

Output size: (19,)

[Hide output](#)

Submit

You have used 2 of 20 attempts

✓ Correct (5/5 points)

## Binary classification error

5/5 points (graded)

Report the test error by running `run_svm_one_vs_rest_on_MNIST`.

Error =  ✓

Submit

You have used 1 of 20 attempts

✓ Correct (5/5 points)

## Implement C-SVM

5/5 points (graded)

Play with the C parameter of SVM, what statement is true about the C parameter?

(Choose all that apply.)

☐ Larger C gives larger tolerance of violation.

☒ Larger C gives smaller tolerance of violation.

☐ Larger C gives a larger-margin separating hyperplane.

☒ Larger C gives a smaller-margin separating hyperplane.

✓

Submit

You have used 1 of 2 attempts

✓ Correct (5/5 points)

## Multiclass SVM

5/5 points (graded)

In fact, `sklearn` already implements a multiclass SVM with a one-vs-rest strategy. Use `LinearSVC` to build a multiclass SVM model

**Available Functions:** You have access to the sklearn's implementation of the linear SVM as `LinearSVC`; No need to import anything.

```
1 def multi_class_svm(train_x, train_y, test_x):
2     """
3     Trains a linear SVM for multiclass classification using a one-vs-rest strategy
4
5     Args:
6         train_x - (n, d) NumPy array (n datapoints each with d features)
7         train_y - (n, ) NumPy array containing the labels (int) for each training data point
8         test_x - (m, d) NumPy array (m datapoints each with d features)
9     Returns:
10         pred_test_y - (m,) NumPy array containing the labels (int) for each test data point
11     """
12     clf = LinearSVC(random_state = 0, C=0.1)
13     clf.fit(train_x, train_y)
14     return clf.predict(test_x)
15
```

Press ESC then TAB or click outside of the code editor to exit

Correct

## Test results

**CORRECT**

[Hide output](#)

Test: output1

Output:



```
train_x: [[6.99578995e-01 5.53887478e-01]
 [6.85779158e-01 6.31515219e-01]
 [7.34893186e-01 3.85731936e-01]
 [8.10086613e-01 6.77818801e-01]
 [4.61409485e-01 8.28452598e-01]
 [4.42838547e-01 9.63062400e-01]
 [2.18127805e-01 2.58773532e-01]
 [5.99256262e-04 9.77156129e-02]
 [9.08444152e-01 2.19172952e-01]
 [6.81866122e-01 4.31294184e-01]
 [5.82768423e-01 2.83390240e-01]
 [7.43304318e-01 2.04510477e-01]
 [8.39572330e-01 1.92496580e-01]]
train_y: [0 1 2 3 3 0 2 3 3 0 2 3 3]
test_x: [[0.54078115 0.90424663]
 [0.30404909 0.72121012]
 [0.9573102 0.54364279]
 [0.78040913 0.4710401 ]
 [0.4121865 0.85725592]
 [0.38522828 0.47252131]
 [0.8350261 0.98659789]
 [0.12303874 0.56828995]
 [0.8581817 0.11108622]
 [0.53666229 0.7750237 ]
 [0.498776 0.30302112]
 [0.0691167 0.80248447]
 [0.81882352 0.71891433]
 [0.11216922 0.30237022]
 [0.95167904 0.58994716]]
Submission output: [3 3 3 3 3 3 3 3 3 3 3 3 3 3 3]
```

Test: output2

Output:

```
train_x: [[0.89067904 0.10903467 0.85626622]
 [0.05025322 0.63078981 0.24687799]
 [0.19626511 0.37969011 0.89536792]
 [0.65222839 0.0262676 0.06054459]
 [0.561349 0.22484606 0.09531271]
 [0.27482608 0.63519925 0.08724291]
 [0.61660784 0.25366854 0.14584416]
 [0.02046448 0.12280837 0.66150661]
 [0.04786739 0.30779211 0.88858349]
 [0.34419749 0.80909343 0.13903339]]
train_y: [3 0 0 0 2 1 0 0 1 2]
test_x: [[0.96478668 0.65586213 0.57584791]
 [0.73845802 0.45179801 0.98143644]
 [0.65385843 0.34829907 0.94578283]
 [0.19782722 0.04934819 0.9415929 ]
 [0.24559923 0.66968244 0.99222757]
 [0.05170463 0.39427076 0.06792789]
 [0.55638174 0.74970882 0.51336906]
 [0.00490819 0.84121466 0.34129819]
 [0.06985683 0.0293824 0.17444237]
 [0.92886513 0.90440363 0.16560619]
 [0.17015272 0.05538263 0.36379799]
 [0.49996439 0.95682557 0.5311442 ]
 [0.00390922 0.71012281 0.83096032]
 [0.32613772 0.18519645 0.73048977]
 [0.00377797 0.67512294 0.26130904]
 [0.93387774 0.02629033 0.08979667]
 [0.87688284 0.71547598 0.82567368]
 [0.46420684 0.99850153 0.22575559]
 [0.96575274 0.31854653 0.62150942]]
Submission output: [0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
```

Test: output3

Output:

```
train_x: [[0.75002148 0.63795031 0.02730299]
[0.06007064 0.92071922 0.95779576]
[0.32853301 0.87291641 0.86004731]
[0.24639947 0.2632307 0.47235945]
[0.05084266 0.76128843 0.94864315]
[0.47737458 0.88383301 0.50558399]
[0.04090603 0.78157897 0.09865894]
[0.75605868 0.75261111 0.4173839 ]
[0.20532656 0.09058549 0.4575146 ]
[0.99151513 0.8016871 0.31725206]
[0.63073877 0.19078004 0.42412705]]
train_y: [1 2 0 3 2 3 4 3 3 3 5]
test_x: [[0.6167434 0.37653891 0.32302156]
[0.52221961 0.09780199 0.8823878 ]
[0.61515609 0.18121164 0.14613885]
[0.59838007 0.50415017 0.54354423]
[0.07601381 0.39909316 0.41933142]
[0.69729208 0.18875622 0.73619069]
[0.09658408 0.68687144 0.54239389]
[0.12313501 0.80736059 0.72992841]
[0.04087764 0.13969736 0.85577583]
[0.36916368 0.25326968 0.46333089]]
Submission output: [3 3 3 3 3 3 3 3 3 3]
```

Test: outputr4

**Output:**

```
train_x: [[0.15113639 0.83737355]
 [0.46977631 0.61996902]
 [0.49968899 0.11361127]
 [0.8626747 0.8358077 ]
 [0.66348748 0.36102256]
 [0.85746564 0.95724941]
 [0.27952378 0.95736125]
 [0.22205355 0.63287309]
 [0.67004944 0.61885252]
 [0.79073745 0.07174858]
 [0.27565991 0.08344264]]
train_y: [0 1 0 2 1 0 3 3 2 0 2]
test_x: [[0.44307091 0.37645626]
 [0.25656735 0.98210213]
 [0.54969574 0.77414383]
 [0.95812598 0.70050533]
 [0.18251712 0.1872745 ]
 [0.95824287 0.98384452]
 [0.78242729 0.06993624]
 [0.96004066 0.62649782]
 [0.48468815 0.79527531]
 [0.14305605 0.55477047]
 [0.93962422 0.16525797]
 [0.14296584 0.46005 ]
 [0.30824337 0.49591007]]
Submission output: [0 0 0 0 0 0 0 0 0 0 0 0]
```

Test: size

Output:

Output size: (12,)

[Hide output](#)

Submit

You have used 1 of 20 attempts

✓ Correct (5/5 points)

Multiclass SVM error

5/5 points (graded)

Report the overall test error by running `run_multiclass_svm_on_MNIST` .

Error = 

0.08189999999999997

 ✓

Submit

 You have used 1 of 20 attempts

✓ Correct (5/5 points)

Discussion

Hide Discussion

**Topic:** Unit 2 Nonlinear Classification, Linear regression, Collaborative Filtering (2 weeks):Project 2: Digit recognition (Part 1) / 3. Support Vector Machine

Add a Post

Show all posts ▼	by recent activity ▼
<div>💬 <a href="#">How to view the misclassified images along with their labels</a></div> <div>To get an idea of what the digit images that were classified incorrectly actually look like, I added a bit of function to display just the misclassified images and put a label above ...</div>	3 ▼
<div>✓ [STAFF] <a href="#">Problem with "Binary classification error"?</a></div>	8 ▼
<div>💬 <a href="#">Implement C-SVM</a></div>	6 ▼
<div>💬 [Staff] <a href="#">Multiclass SVM</a></div> <div>Grader erroneously asks to implement `one vs rest svm` twice instead of `multi class svm`.</div> <div>👤 <a href="#">Community TA</a></div>	11 ▼
<div>💬 <a href="#">Error in Answer: Implement C-SVM</a></div>	7 ▼
<div>✓ <a href="#">Multiclass SVM error</a></div> <div>Multiclass SVM works fine and passes the grader, but the error -- which is lower than for one vs. rest SVM -- is marked as wrong.</div>	6 ▼

Learn About Verified Certificates