edX

MITx: 6.86x
Machine Learning with Python-From Linear Models to Deep Learning

Help   sandipan_dey

4. RNN Decoding
Decoding
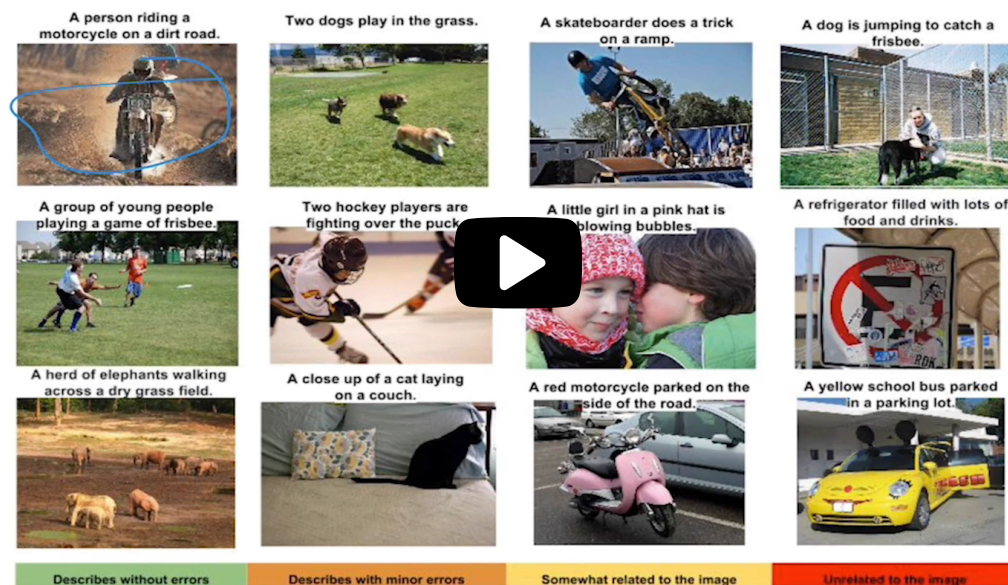
So some things are generated without an error.

Here is an image.

And what the model predicts corresponding to this image

is there a person riding a motorcycle on a dirt road,

so very accurate.

And it makes also some errors.

The thing to understand here is really

that since these are very complicated models,

they can also memorize the examples quite well.

So in order to understand how well they actually do,

you would have to compare these sentences to the sentences that

were available during training.

So it's not just picking your nearest neighbor

of those sentences to associate with the image.

So key things to understand from this lecture

▶ 2:29 / 3:44   ▶ Speed 1.50x   🔊   ⛶   CC   ❝

Video
Download video file

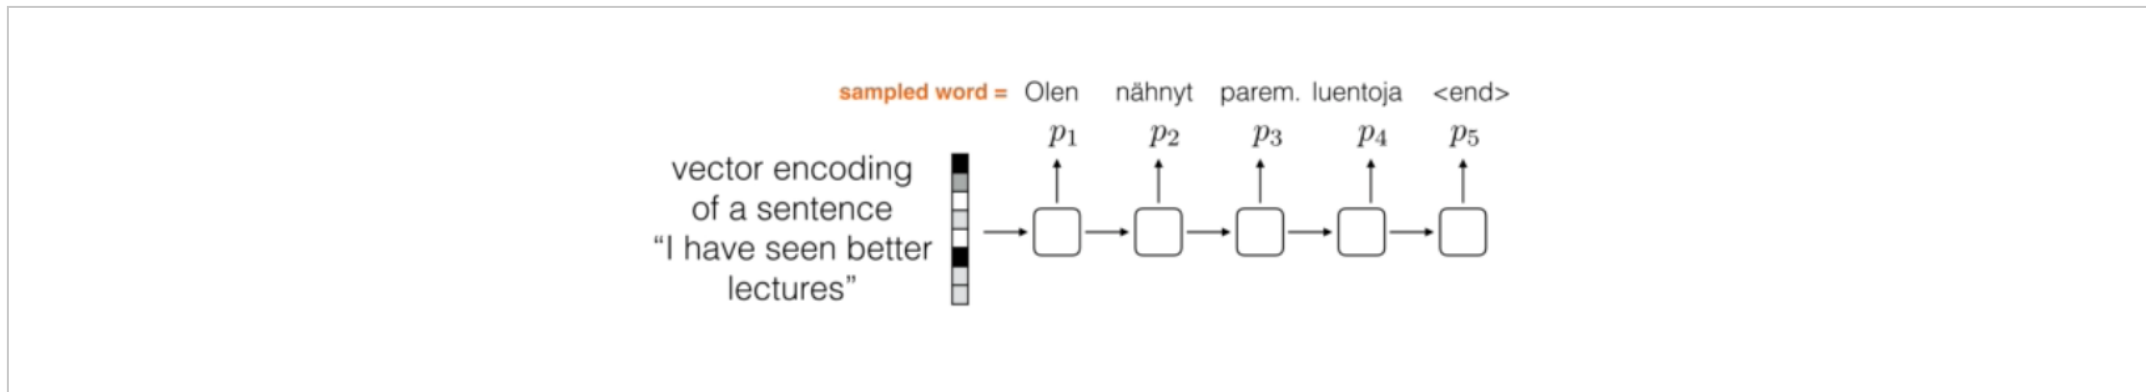## Transcripts
Download SubRip (.srt) file
Download Text (.txt) file

Decoding with RNN

1/1 point (graded)

Now, we would like to decode a feature vector with RNN's. The picture below illustrates how a vector encoding of the English sentence "I have seen better lectures" is translated into a sentence of a foreign language.

Unlike in encoding, at each step, an output distribution p_t is produced in a decoding RNN.

$$s_t = \tanh(W^{s,s} s_{t-1} + W^{s,x} x_t) \quad \text{state}$$
$$p_t = \text{softmax}(W^o s_t) \quad \text{output distribution}$$
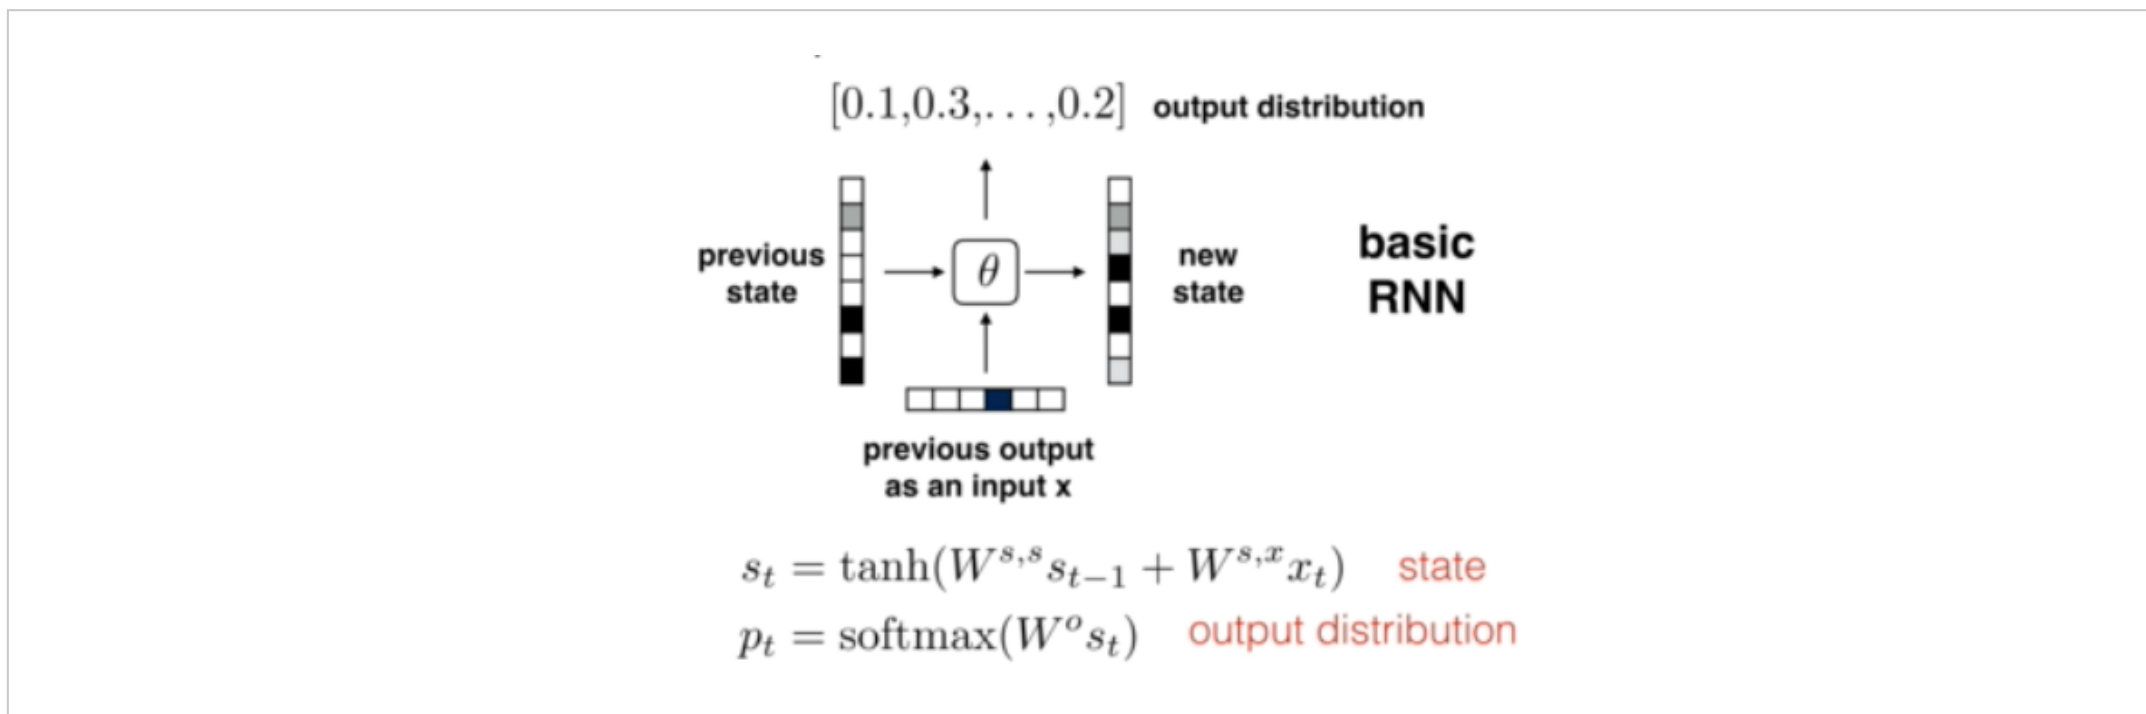
Now, which of the following is true about decoding RNN's? Choose all those apply.

☐ In the translating example above, the output probability distribution p_t is fed as an input to the next step

☐ The probability distribution p_t is the same at each step, just like how parameters are shared between steps

☑ In the first image, the foreign word "Olen" in the above picture is a "sampled" result from a distribution the RNN produced. ✔

✔

Solution:

As shown in the figure, it is the previous output x_t but not the output probability ditribution p_t is fed into the next step. The probability distribution is different at each step as it propogated from the begining state. With the probability distribution at each step, the output word is then sampled from the distribution.

Submit	You have used 1 of 2 attempts

ⓘ   Answers are displayed within the problem

Predictions

1/1 point (graded)

Suppose we are building an RNN model to translate images into sentences, as described in the lecture. Which of the following is only done in generating predictions from an RNN on test data?

☑ Feeding the sampled output as part of the input to the next time step ✔

☐ Calculating what percentage of words the RNN correctly generated

☐ Feeding in the hidden state as input each time step

✔

Solution:

All of the above are done during training except using the sampled output as input. In training, you use the true words specified as input for the next time step. However, in testing you want to predict the sentence on your own, so you use the sampled output at one time step as the input for the next step.

Submit	You have used 2 of 2 attempts

ⓘ   Answers are displayed within the problem

Discussion

Topic: Unit 3 Neural networks (2.5 weeks):Lecture 11. Recurrent Neural Networks 2 / 4. RNN Decoding

Add a Post

Show all posts   ▼   by recent activity ▼

💬   Encoder-decoder rnn

For an intuitive description of rnn encoding/decoding: https://devblogs.nvidia.com/introduction-neural-machine-translation-with-gpus/ https://devblogs.nvidia.com/introducti...   1

👤 Community TA

Learn About Verified Certificates