EdX and its Members use cookies and other tracking technologies for performance, analytics, and marketing purposes. By using this website, you accept this use. Learn more about these technologies in the <u>Privacy Policy</u>.





Course > Unit 3 Neural networks (2.5 weeks) > Project 3: Digit recognition (Part 2) > 9. Convolutional Neural Networks

## 9. Convolutional Neural Networks

Next, we are going to apply convolutional neural networks to the same task. These networks have demonstrated great performance on many deep learning tasks, especially in computer vision.

You will be working in the files part2-mnist/nnet\_cnn.py and part2-mnist/train\_utils.py in this problem

### Convolutional Neural Networks

3/3 points (graded)

We provide skeleton code <code>part2-mnist/nnet\_cnn.py</code> which includes examples of some (**not all**) of the new layers you will need in this part. Using the <u>PyTorch Documentation</u>, complete the code to implement a convolutional neural network with following layers in order:

- $\bullet\,$  A convolutional layer with 32 filters of size  $3\times3$
- A ReLU nonlinearity

- A max pooling layer with size  $2 \times 2$
- ullet A convolutional layer with 64 filters of size 3 imes 3
- A ReLU nonlinearity
- ullet A max pooling layer with size 2 imes 2
- A flatten layer
- A fully connected layer with 128 neurons
- A dropout layer with drop probability 0.5
- A fully-connected layer with 10 neurons

**Note:** We are not using a softmax layer because it is already present in the loss: PyTorch's <a href="mailto:nn.CrossEntropyLoss">nn.CrossEntropyLoss</a> combines <a href="mailto:nn.NLLLoss">nn.NLLLoss</a>.

Without GPU acceleration, you will likely find that this network takes quite a long time to train. For that reason, we don't expect you to actually train this network until convergence. Implementing the layers and verifying that you get approximately 93% **training accuracy** and 98% **validation accuracy** after one training epoch (this should take less than 10 minutes) is enough for this project. If you are curious, you can let the model train longer; if implemented correctly, your model should achieve >99% **test accuracy** after 10 epochs of training. If you have access to a CUDA compatible GPU, you could even try configuring PyTorch to use your GPU.

After you successfully implement the above architecture, copy+paste your model code into the codebox below for grading.

**Available Functions:** You have access to the torch.nn module as nn and to the Flatten layer as Flatten; No need to import anything.

```
1 model = nn.Sequential(
 2
            nn.Conv2d(1, 32, (3, 3)),
 3
             nn.ReLU(),
             nn.MaxPool2d((2, 2)),
 4
 5
             nn.Conv2d(32, 64, (3, 3)),
 6
             nn.ReLU(),
            nn.MaxPool2d((2, 2)),
 8
             Flatten(),
 9
             nn.Linear(1600, 128),
            nn.Dropout(0.5),
10
11
            nn.Linear(128, 10),
12
```

Press ESC then TAB or click outside of the code editor to exit

Correct

# Test results

```
CORRECT

Test: has 11 layers

Testing has 11 layers

Output:
```

Length: 10 Test: has correct layers Testing has correct layers Output: Layers: Conv2d ReLU MaxPool2d Conv2d ReLU MaxPool2d Flatten Linear Dropout Linear Test: has correct number of parameters Testing has correct number of parameters Output: Number of parameters: 225034

Test: has correct output

Testing has correct output

Output:

```
model(a): tensor([[ 2.4448, 1.5225, 2.1802, 4.4122, 0.5503, 1.3307, -2.1938, 2.3987,
        -2.9601, -0.4598]])
a[:,:,:5] tensor([[[ 1.1542e-01, 1.3941e+00, -2.7714e-01, 4.2701e-01, 9.3229e-01,
           2.3895e-01, 3.2457e-01, -3.4939e-01, 1.9670e-02, 2.2480e-01,
           1.9173e+00, 2.2247e-01, -4.6467e-01, -7.6888e-01, 3.0430e-02,
           9.3643e-01, 7.7472e-02, 3.7119e-01, -1.7743e+00, -7.9063e-01,
           3.8385e-02, 6.5397e-01, -8.9731e-02, -1.4515e-01, 3.5549e-01,
          -7.8692e-01, -5.8899e-04, -4.8655e-01],
         [ 1.3405e+00, 1.2712e+00, 3.7503e-01, 8.9544e-01, -4.0794e-01,
           2.1923e+00, -6.6135e-01, -2.2415e+00, -1.4822e+00, -5.5531e-01,
           8.7844e-01, 8.4059e-01, 5.6392e-01, 1.8175e-01, 5.1660e-01,
           3.1636e-01, 7.3709e-01, 2.5617e-02, -2.5206e-01, -4.5273e-01,
           1.5415e+00, -1.5221e+00, -7.3029e-01, 2.2291e-01, 4.1424e-01,
          -3.2722e-01. -1.5834e-01. -4.2584e-011.
         [-5.0879e-01, -6.4091e-01, -2.4095e-01, -2.6180e-01, 9.3761e-01,
          -1.3980e-02, -7.2690e-01, -8.9905e-02, 2.6777e-01, 1.2556e-01,
           1.4992e+00, -8.0736e-01, 3.7684e-01, 2.7557e-01, 1.7586e-01,
          -7.9403e-01, 4.4957e-01, 7.0530e-02, 1.3219e-01, 5.7742e-01,
           1.0787e-01, -1.1266e+00, -3.7247e-01, -1.6174e+00, 6.5828e-01,
          -7.1602e-01, -6.8528e-01, 6.9362e-01],
         [-1.2379e+00, 2.4469e-01, 2.4287e-01, -2.3104e+00, -1.9266e-02,
           4.6654e-01, 6.2467e-01, -1.1184e+00, 8.1859e-01, 6.2088e-01,
           6.2070e-01. 1.7712e-01. -3.1773e-01. -6.1871e-01. -5.3287e-01.
           1.0487e+00, 1.8772e+00, 9.3554e-01, -1.2786e-02, -7.4107e-02,
           1.0031e+00, -9.3705e-01, -1.1818e+00, 8.6155e-01, -1.1396e+00,
           3.6708e-01, -7.9953e-01, -7.9232e-01],
         [ 3.1212e-02, -7.6052e-01, -3.5220e-01, 1.0426e-01, 8.0699e-01,
           9.4311e-01, 1.9873e+00, 1.5439e+00, -9.5529e-02, -1.6095e+00,
          -6.9276e-01, -7.7082e-01, 4.8388e-02, -8.1956e-02, 1.2461e-02,
          -1.0657e+00, 8.7008e-01, -1.5414e+00, -1.3320e+00, -6.0466e-01,
          -1.1799e-01, -5.1236e-01, 1.5981e-01, 7.8021e-01, -2.4048e-01,
          -9.0496e-01, -5.5437e-01, -1.3648e+001111)
```

**Hide output** 

Submit

You have used 7 of 20 attempts

✓ Correct (3/3 points)

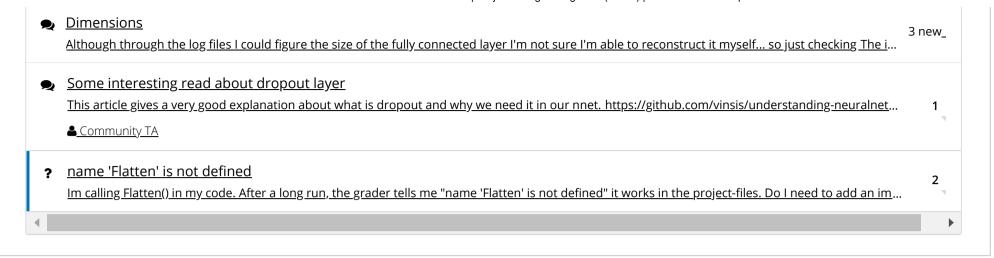
## Discussion

**Hide Discussion** 

**Topic:** Unit 3 Neural networks (2.5 weeks):Project 3: Digit recognition (Part 2) / 9. Convolutional Neural Networks

#### Add a Post

Show all posts by recent activity ▼ • [Staff] There is a grader problem (NameError: name 'Flatten' is not defined). 1 new Dear Staff, It seems something wrong with `Flatten` definition in the grader. I've got the following problem twice: Traceback (most recent call last): Fi... Grader processing too long for 9. Convolutional Neural Networks 23 I submitted the code, and the grader has already been running for 1 hour, .. and I cannot submit new version of code... is that normal? Useful link for those struggling like me 4 I found the following link useful https://blog.algorithmia.com/convolutional-neural-nets-in-pytorch/ flatten and fully connected layers 3 1. If I understand correctly, fully connected layer is a Linear layer. We are given only one dimension of it (128). What is the other param? Is it 10 (sa... what model code needs to be pasted? 6 Lcreated a class for my model as described in this article [Convolutional Neural Networks Tutorial in PyTorch ][1] <br/>
Show do I post the code into gr...



© All Rights Reserved