

EdX and its Members use cookies and other tracking technologies for performance, analytics, and marketing purposes. By using this website, you accept this use. Learn more about these technologies in the [Privacy Policy](#).



[Course](#) > [Unit 3 Neural networks \(2.5 weeks\)](#) > [Project 3: Digit recognition \(Part 2\)](#) > 10. Overlapping, multi-digit MNIST

10. Overlapping, multi-digit MNIST

In this problem, we are going to go beyond the basic MNIST. We will train a few neural networks to solve the problem of hand-written digit recognition using a multi-digit version of MNIST.



You will be working in the files `part2-twodigit/mlp.py`, `part2-twodigit/conv.py`, and `part2-twodigit/train_utils.py` in this problem

In your project folder, look at the **part2-twodigit** subfolder. There you can find the files **mlp.py** and **conv.py**. Your main task here is to complete the code inside the method `main` in these files.

Do the following steps:

- Look at `main` method in each file. Identify the training and test data and labels. How many images are inside the train and test data? What is the size of each image?
- Look at the definition of the MLP class in **mlp.py**. Try to make sense of what those lines are trying to achieve. What is `y_train[0]` and `y_train[1]`?
- Look at `train_utils.py`, particularly the `run_epoch` function.

Now given the intuition you have built with the above steps, complete the following tasks.

Fully connected network

5.0/5.0 points (graded)

Complete the code **main** in **mlp.py** to build a fully-connected model with a single hidden layer with 64 units. For this, you need to make use of `Linear` layers in PyTorch; we provide you with an implementation of `Flatten`, which maps a higher dimensional tensor into an $N \times d$ one, where N is the number of samples in your batch and d is the length of the flattened dimension (if your tensor is $N \times h \times w$, the flattened dimension is $d = (h \cdot w)$). Hint: Note that your model must have two outputs (corresponding to the first and second digits) to be compatible with the data.

Available Functions: You have access to the `torch.nn` module as `nn`, to the `torch.nn.functional` as `F` and to the `Flatten` layer as `Flatten`; No need to import anything.

```
1 class MLP(nn.Module):
2
3     def __init__(self, input_dimension):
4         super(MLP, self).__init__()
5         self.flatten = Flatten()
6         self.fc1 = nn.Sequential(nn.Linear(input_dimension, 64),
7                                   nn.ReLU(),
8                                   nn.Linear(64, 10),)
9         self.fc2 = nn.Sequential(nn.Linear(input_dimension, 64),
10                                   nn.ReLU(),
11                                   nn.Linear(64, 10),)
12
13     def forward(self, x):
14         xf = self.flatten(x)
15         # ... (rest of the code) ...
```

Press ESC then TAB or click outside of the code editor to exit

Correct

Test results

[Hide output](#)

CORRECT

Test: can overfit simpledata

Testing can overfit simpledata

Output:

```
Final loss < 0.5: True
```

Test: has correct io

```
Testing has init method
```

Output:

```
Size of out1: torch.Size([2, 10])  
Size of out2: torch.Size([2, 10])
```

Test: has init method

```
Testing has init method
```

Output:

```
Initialized
```

[Hide output](#)

Submit

You have used 4 of 20 attempts

Convolutional model

5.0/5.0 points (graded)

Complete the code `main` in **conv.py** to build a convolutional model. For this, you need to make use of **Conv2D** layers and **MaxPool2d** layers (and perhaps Dropout) in PyTorch. Make sure that the last layer of the neural network is a fully connected (Linear) layer.

Available Functions: You have access to the `torch.nn` module as `nn`, to the `torch.nn.functional` as `F` and to the `Flatten` layer as `Flatten`; No need to import anything.

```
1 class CNN(nn.Module):
2
3     def __init__(self, input_dimension):
4         super(CNN, self).__init__()
5         # TODO initialize model layers here
6         self.model1 = nn.Sequential(
7             nn.Conv2d(1, 32, (3, 3)),
8             nn.ReLU(),
9             nn.MaxPool2d((2, 2)),
10            nn.Conv2d(32, 64, (3, 3)),
11            nn.ReLU(),
12            nn.MaxPool2d((2, 2)),
13            Flatten(),
14            nn.Linear(2880, 128),
15            nn.Dropout(0.5),
16            nn.ReLU()
```

Press ESC then TAB or click outside of the code editor to exit

Correct

Test results

[Hide output](#)**Correct:**

Test: has correct layers

Testing has correct layers

Output:

```
Conv2d
MaxPool2d
Dropout
```

Test: has init method

Testing has init method

Output:

```
Initialized
```

[Hide output](#)**Submit**

You have used 3 of 20 attempts

Hyperparameter tuning

1/1 point (graded)

Next, change the parameters and settings of the models above. Train your model with various parameter settings. Some things you can try is to modify the learning algorithm from **SGD** to more sophisticated ones (such as **ADAM**) or you can modify the network architecture (number of layers or unit per layers, activation function, etc.). Something to ponder: What parameters or settings were more conducive to get a better model with greater generalization capability and lower error? Did extra training for a model always help or did the training accuracy plateau or even get worse after some point?

Finally, we will grade you on finding at least one architecture that achieves over 98% accuracy on both the validation and test set.

Please enter your **test accuracy** .

0.980595



Submit

You have used 1 of 5 attempts

Conclusion and What's Next

As you have seen in this project, neural networks can pretty successfully solve the MNIST task. In fact, since 2012, following the impressive performance of AlexNet on the ImageNet dataset, deep neural networks have been the standard in computer vision. As datasets went growing in size and complexity and as computing power became cheaper and more efficient, the trend has been to build deeper and bigger neural nets.

The last part of the project has given you a hint as to why neural networks can be very versatile: by merely changing the output layer, you were able to train the network to predict overlapping MNIST digits. The same building blocks can be reused to build more complex architecture and solve more difficult problems. Using a deep learning framework like Pytorch makes

this process even more accessible.

If you have access to a GPU, you can try implementing an object classification system with Resnet, which we have not covered in this course, and maybe expanding it to an object detection or an image segmentation system.

If you do not have access to a GPU, you can try renting resources from an online provider, such as [Paperspace](#) (<1\$/hour) or [Google Colab](#) (free).

Discussion

Topic: Unit 3 Neural networks (2.5 weeks);Project 3: Digit recognition (Part 2) / 10. Overlapping, multi-digit MNIST

Hide Discussion

Add a Post

Show all posts ▼		by recent activity ▼	
?	Architectures of MLP and CNN ----- edited out Mark B2 -----	5	
💬	num_classes not defined anywhere Although the file mlp.py contains a definition of the variable, the grader testing code does not have it defined anywhere, although i think it's exp... ★ Following	1	
?	Hyperparameter tuning How does the grader check the answer ?	1	
?	What is forward ? What is MLP ? I m lost here. What is the first step. I don't understand what is asked here.	2	

?	[STAFF] Please Help !!!!!	10
	I watched the course carefully. I went trough the documentation of pyTorch. I have no idea what I have to do and how to do it. Please provide us...	
💬	Grader is down	9
	Processing for 24 hrs... Please take a look	
💬	A lazy way to implement pytorch GPU support	1
	For whom interested in grid search :) My aim is minimum code change and avoid disrupting exiting template. It's not most efficiency method, bu...	
	Community TA	
💬	Classifying ordered pairs or each digit individually?	3
	For the last part I found parameters with 98.5 percent accuracy in the first digit and 97.9 percent accuracy in the 2nd digit which I'd like to count ...	
?	Any hint on how to predict 2 digits instead of 1 digit	9
	When we want to predict one digit, we construct a final dense layer with output 10, but how about now we are predict 2 digits, how should we co...	
?	nn.Softmax vs. nn.functional.softmax	4
	What is/are the difference(s) between the two ways of implementing softmax ?	
?	failing 'can overfit simple data'	6
	I keep failing this test for the MLP class. I'm producing about 90 percent accuracy on test set locally. I didn't see much freedom in the instruction...	
?	Any clue from the teaching staff, mine to help others is conv nn, split, split ... divide and conquer.	2
	This is not good:
 Epoch 30:
 100% ██████████ 562/562 [01:28<00:00, 6.11it/s]
 Train loss1: 0.045951 accuracy1: 0.983736 lo...	

Learn About Verified Certificates