





Branch: master ▾ pytorch-tutorial / tutorials / 02-intermediate / convolutional_neural_network / main.py Find file Copy path

 yunjey Update tutorials for pytorch 0.4.0 78c6afe on May 10, 2018

1 contributor

100 lines (82 sloc) 3.28 KB Raw Blame History   

```
1 import torch
2 import torch.nn as nn
3 import torchvision
4 import torchvision.transforms as transforms
5
6
7 # Device configuration
8 device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
9
10 # Hyper parameters
11 num_epochs = 5
12 num_classes = 10
13 batch_size = 100
14 learning_rate = 0.001
15
16 # MNIST dataset
17 train_dataset = torchvision.datasets.MNIST(root='../data/',
18                                           train=True,
19                                           transform=transforms.ToTensor(),
20                                           download=True)
21
22 test_dataset = torchvision.datasets.MNIST(root='../data/',
23                                           train=False,
24                                           transform=transforms.ToTensor())
25
26 # Data loader
27 train_loader = torch.utils.data.DataLoader(dataset=train_dataset,
28                                             batch_size=batch_size,
29                                             shuffle=True)
30
31 test_loader = torch.utils.data.DataLoader(dataset=test_dataset,
32                                           batch_size=batch_size,
33                                           shuffle=False)
34
35 # Convolutional neural network (two convolutional layers)
36 class ConvNet(nn.Module):
37     def __init__(self, num_classes=10):
38         super(ConvNet, self).__init__()
39         self.layer1 = nn.Sequential(
```

```
40         nn.Conv2d(1, 16, kernel_size=5, stride=1, padding=2),
41         nn.BatchNorm2d(16),
42         nn.ReLU(),
43         nn.MaxPool2d(kernel_size=2, stride=2))
44     self.layer2 = nn.Sequential(
45         nn.Conv2d(16, 32, kernel_size=5, stride=1, padding=2),
46         nn.BatchNorm2d(32),
47         nn.ReLU(),
48         nn.MaxPool2d(kernel_size=2, stride=2))
49     self.fc = nn.Linear(7*7*32, num_classes)
50
51     def forward(self, x):
52         out = self.layer1(x)
53         out = self.layer2(out)
54         out = out.reshape(out.size(0), -1)
55         out = self.fc(out)
56         return out
57
58 model = ConvNet(num_classes).to(device)
59
60 # Loss and optimizer
61 criterion = nn.CrossEntropyLoss()
62 optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
63
64 # Train the model
65 total_step = len(train_loader)
66 for epoch in range(num_epochs):
67     for i, (images, labels) in enumerate(train_loader):
68         images = images.to(device)
69         labels = labels.to(device)
70
71         # Forward pass
72         outputs = model(images)
73         loss = criterion(outputs, labels)
74
75         # Backward and optimize
76         optimizer.zero_grad()
77         loss.backward()
78         optimizer.step()
79
80         if (i+1) % 100 == 0:
81             print ('Epoch [{}/{}], Step [{}/{}], Loss: {:.4f}'
82                   .format(epoch+1, num_epochs, i+1, total_step, loss.item()))
83
84 # Test the model
85 model.eval() # eval mode (batchnorm uses moving mean/variance instead of mini-batch mean/variance)
86 with torch.no_grad():
87     correct = 0
88     total = 0
89     for images, labels in test_loader:
90         images = images.to(device)
91         labels = labels.to(device)
```

```
92         outputs = model(images)
93         _, predicted = torch.max(outputs.data, 1)
94         total += labels.size(0)
95         correct += (predicted == labels).sum().item()
96
97     print('Test Accuracy of the model on the 10000 test images: {} %'.format(100 * correct / total))
98
99     # Save the model checkpoint
100    torch.save(model.state_dict(), 'model.ckpt')
```