



Why do we use a one-tailed test F-test in analysis of variance (ANOVA)?

Can you give the reason for using a one tailed test in the analysis of variance test?

Why do we use a one-tail test - the F-test - in ANOVA?

anova | f-test | sums-of-squares | f-distribution

edited Nov 23 '15 at 10:32



amoeba

30k 8 107 172

asked Aug 16 '13 at 6:36



Cynderella

90 1 1 6

- 1 Some questions to guide your thinking... What does a very negative t statistic mean? Is a negative F statistic possible? What does a very low F statistic mean? What does a high F statistic mean? – [rpierce](#) Aug 16 '13 at 6:41

Why are you under the impression that a one-tailed test has to be an F-Test? To answer your question: The F-Test allows to test a hypothesis with more than one linear combination of parameters. – [IMA](#) Aug 16 '13 at 7:17

- 1 Do you want to know why one would use a one-tailed instead of a two-tailed test? – [Jens Kouros](#) Aug 16 '13 at 7:53

@tree what constitutes a credible or official source for your purposes? – [Glen_b](#) ♦ Oct 27 '13 at 23:24

- 1 @tree note that Cynderella's question here is *not* about a test of variances, but specifically an F-test of ANOVA - which is a test for *equality of means*. If you're interested in tests of equality of variances, that's been discussed in many other questions on this site. (For the variance test, yes, you do care about both tails, as is clearly explained in the last sentence of [this section](#), right above **'Properties'**) – [Glen_b](#) ♦ Oct 28 '13 at 22:25

3 Answers

F tests are most commonly used for two purposes:

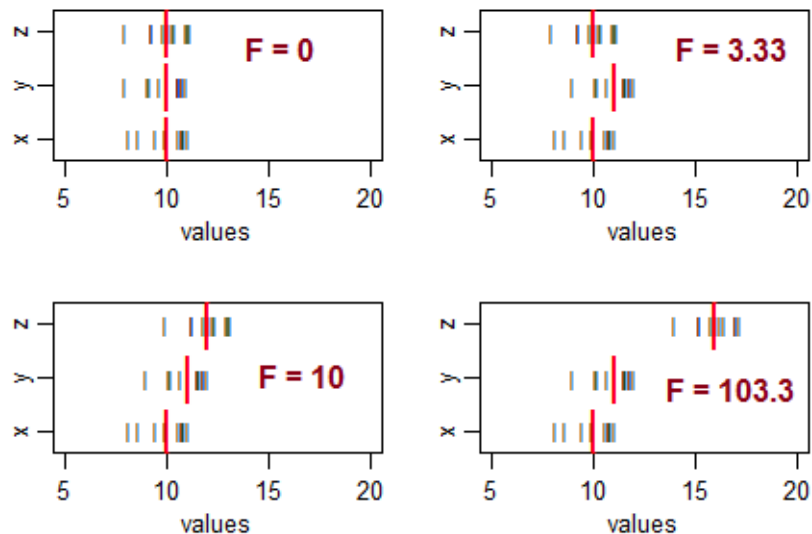
1. in ANOVA, for testing equality of means (and various similar analyses); and

2. in testing equality of variances

Let's consider each in turn:

1) F tests in ANOVA (and similarly, the usual kinds of chi-square tests for count data) are constructed so that the more the data are consistent with the alternative hypothesis, the larger the test statistic tends to be, while arrangements of sample data that looks most consistent with the null corresponds to the smallest values of the test statistic.

Consider three samples (of size 10, with equal sample variance), and arrange them to have equal sample means, and then move their means around in different patterns. As the variation in the sample means increases from zero, the F statistic becomes larger:




The black lines (|) are the data values. The heavy red lines (|) are the group means.

If the null hypothesis (equality of population means) were true, you'd expect some variation in sample means, and would typically expect to see F ratios roughly around 1. Smaller F statistics result from samples that are closer together than you'd typically expect ... so you aren't going to conclude the population means differ.

That is, for ANOVA, you'll reject the hypothesis of equality of means when you get unusually large F-values and you won't reject the hypothesis of equality of means when you get unusually small values (it may indicate *something*, but not that the population means differ).

Here's an illustration that might help you see that we only want to reject when F is in its upper tail:



2) F tests for equality of variance* (based on variance ratios). Here, the ratio of two sample variance estimates will be large if the numerator sample variance is much larger than the variance in the denominator, and the ratio will be small if the denominator sample variance is much larger than variance in the numerator.

That is, for testing whether the ratio of population variances differs from 1, you'll want to reject the null for both large and small values of F.

* (Leaving aside the issue of the high sensitivity to the distributional assumption of this test (there are better alternatives) and also the issue that if you're interested in suitability of ANOVA equal-variance assumptions, your best strategy probably isn't a formal test.)

edited Nov 23 '15 at 4:04

answered Oct 28 '13 at 23:09



Glen_b ♦

153k 20 254 525

2 @TaylerJones Levene's test is somewhat more robust. Browne-Forsythe is more robust (but loses a little power near the normal). Fligner-Killeen more so again. In several decades, I've used Levene or Browne-Forsythe no more than twice each. (If it came up again, likely something like Browne-Forsythe would suit me fine but I don't

generally have situations where it makes any sense to test several groups variances for equality.) – Glen_b ♦ Oct 29 '13 at 7:34

2 I appology. I still haven't understood why do we use one-tail in ANOVA . More specifically, from your discussion i understood that under null hypothesis i would not have any treatment effect and consequently $F = \frac{MS_{TREATMENT}}{MS_{ERROR}}$ will be close to 1 , while if alternative hypothesis is true the F -Ratio will be larger. But how does it imply "*that is the reason for using one-tail test in ANOVA?*" – time Oct 29 '13 at 12:18

2 @tree it sounds like you don't understand something about hypothesis testing more generally, but it's hard to be certain exactly where. You say you understand that if you get a large F you want to reject and if you get a small F you don't want to reject. The large values of F are those values in the upper tail while small values of F are those values in the lower tail. You only want to reject when the values are large ... i.e. in the upper tail, but not the lower tail. How can you not see that's one tailed? I'll include another plot that might help. – Glen_b ♦ Oct 29 '13 at 22:03

1 @jeramy My comments refer to the tests that rely on ratios of variances (specifically, I stated "*Here, the ratio of two sample variance estimates will ...*"). The tests you refer to look for location differences in absolute residuals from some location measure in order to spot differences in spread; they naturally work the way tests for location-differences work. Since I was trying to show a case where you *would* look at the lower tail of the F, the Brown-Forsythe (& some other tests that look for location differences in some measure of deviation to infer spread differences) would be no help – Glen_b ♦ Nov 23 '15 at 4:03

1 @jeramy I have added a few words to make it more explicit. You may like to note that even though Brown-Forsythe, Levene and so on *use* F-tables, the distribution of the test statistics are not actually F-distributed, even under the assumptions of the test. – Glen_b ♦ Nov 23 '15 at 4:05

|

It must be understood that the objective of ANOVA is to check whether there is inequality of means...which implies that we are concerned with large variations between samples (& thus means as variations are computed from the means) as compared with variations within samples (again computed from individual sample mean). When the variations between samples is small (resulting in F value being on the left side) it does not matter as this difference is insignificant. The variations between samples matters if it is significantly higher than the within variations & in such case the F value would be greater than 1, & therefore in the right tail.

The only question remains is why put the entire level of significance in the right tail & the answer is again similar. The rejection happens only when the F ratio is on the right side & never when the F ratio is on the left side. The level of significance is the measure of error due to statistical limitations. As the rejection happens only on the right the entire level of significance (error risk of misconception) is kept in the right. `

edited Oct 1 '14 at 11:35

answered Oct 1 '14 at 10:16



Glen_b ♦

153k 20 254 525



Prof Pradeep Pai

21 1

The expected value for the Mean Square (MS) within treatments is the population variance, whereas the expected value for the MS between treatments is the population variance PLUS the treatment variance. Thus, the ratio of $F = MS_{\text{between}} / MS_{\text{within}}$ is always greater than 1, and never less than 1.

Since the precision of a 1-tailed test is better than a 2-tailed test, we prefer to use the 1-tailed test.

answered Nov 5 '15 at 23:17



Jeff Cotter

1

I don't believe the claim in the last sentence of your first paragraph is correct... $E(\text{numerator}) > E(\text{denominator})$ does not imply that $\text{numerator} > \text{denominator}$. – Glen_b ♦ Nov 6 '15 at 0:17

Aside from Glen_b's point, I'm not sure about "since the precision of a 1-tailed test is better than a 2-tailed test, we prefer to use the 1-tailed test." Can you explain what you mean by this? Talking about precision seems to me to miss the point. – Silverfish Nov 6 '15 at 0:19

Precision is same as half confidence interval. For the same F-stat, a 1 tail test will reject the null hypothesis with a smaller p-value (half, in fact). The other way around, a 1 tail test can reject the null hypothesis with smaller values of the F-stat. This implies that a 1 tail test can detect a treatment effect with fewer samples or with more common cause variance present in the sample. This makes the 1 tail test more desirable, if one is looking for an effect. – Jeff Cotter Nov 6 '15 at 1:23

Yes, a calculated F statistic can be less than 1.0. However, the conclusion would be fail to reject the null hypothesis of "no treatment effects". Therefore, there's no critical region in the lower tail. Therefore, the F-test is an upper one-tailed test. In ANOVA, the logical argument is based on the expected values for MS_{treat} and MS_{error} . Under the "no treatment effect" hypothesis, $H_0: E(MS_{\text{treat}}) = E(MS_{\text{error}}) = \text{population variance}$. Any significant treatment effect results in $H_A: E(MS_{\text{treat}}) > E(MS_{\text{error}})$. (Source any Montgomery text covering ANOVA). Thus, H_A implies a one-tailed test. – Jeff Cotter Nov 6 '15 at 1:30