


Learn Git and GitHub without any code!

Using the Hello World guide, you'll start a branch, write comments, and open a pull request.

[Read the guide](#)

Branch: master ▾ [bayesian-deep-learning](#) / [examples](#) / [notebooks](#) / [Anomaly_Detection_via_VAEs.ipynb](#) [Find file](#) [Copy path](#)

 **seungjin-yang** Add an example about anomaly detection via VAE b0cca73 on Oct 18

1 contributor

930 lines (929 sloc) 132 KB ⏏ 📄 [Raw](#) [Blame](#) [History](#) 🖨 ✎ 🗑

Start with Google Colab (<https://colab.research.google.com/github/seungjin-yang/bayesian-deep-learning/blob/master/examples/notebooks/Anomaly-Detection-via-VAEs.ipynb>)

based on TFP Probabilistic Layers VAE (https://github.com/tensorflow/probability/blob/master/tensorflow_probability/examples/jupyter_notebooks/Probabilistic_Layers_VAE.ipynb)

```
In [ ]: try:
        import google.colab
        IS_COLAB = True
    except:
        IS_COLAB = False

    if IS_COLAB:
        ! pip install tensorflow-gpu==2.0.0
        ! pip install tensorflow-probability==0.8.0
        ! pip install tensorflow-datasets==1.2.0
        ! pip install git+https://github.com/seungjin-yang/bayesian-deep-learning
```

```
In [1]: %matplotlib inline
import numpy as np
import matplotlib.pyplot as plt

import tensorflow as tf
import tensorflow_probability as tfp
import tensorflow_datasets as tfds

import extended_tfp as extfp

tfk = tf.keras
tfkl = tf.keras.layers
tfpl = tfp.layers
tfd = tfp.distributions

tf.keras.backend.set_image_data_format('channels_first')
```

```
In [2]: print(f'TensorFlow: {tf.version.VERSION}')
print(f'TensorFlow Probbility: {tfp.__version__}')
print(f'TensoroFlow Datasets: {tfds.__version__}')

TensorFlow: 2.0.0
TensorFlow Probbility: 0.8.0
TensoroFlow Datasets: 1.2.0
```

Data

```
In [3]: datasets, datasets_info = tfds.load(name='mnist', with_info=True, as_supervised=False)
```

WARNING: Logging before flag parsing goes to stderr

bayesian-deep-learning/Anomaly_Detection_via_VAEs.ipynb at master · seungjin-yang/bayesian-deep-learning
 WARNING: Logging before flag parsing goes to stderr.
 W1010 11:42:27.887978 140359066490688 dataset_builder.py:439] Warning: Setting shuffle_files=True because split=TRAIN and shuffle_files=None. This behavior will be deprecated on 2019-08-06, at which point shuffle_files=False will be the default for all splits.

```
In [4]: def preprocess_mnist(sample):
        # Scale to unit interval.
        image = tf.cast(sample['image'], tf.float32) / 255.
        # NHWC to NCHW
        image = tf.transpose(image, (2, 0, 1))

        # Randomly binarize.
        # image = image < tf.random.uniform(tf.shape(image))

        input_image = tf.cast(image, tf.float32)
        label_image = tf.cast(image, tf.int64)
        return input_image, image

train_dataset = (datasets['train']
                 .map(preprocess_mnist)
                 .batch(256)
                 .prefetch(tf.data.experimental.AUTOTUNE)
                 .shuffle(int(10e3)))

eval_dataset = (datasets['test']
                .map(preprocess_mnist)
                .batch(256)
                .prefetch(tf.data.experimental.AUTOTUNE))
```

```
In [7]: def build_model(input_shape, encoded_size=16, base_depth=32):
        prior = tfd.Independent(
            tfd.Normal(loc=tf.zeros(encoded_size), scale=1),
            reinterpreted_batch_ndims=1)

        # posterior
        encoder = tfk.Sequential([
            tfkl.InputLayer(input_shape=input_shape),
            tfkl.Conv2D(base_depth, 5, strides=1,
                        padding='same', activation=tf.nn.leaky_relu),
            tfkl.Conv2D(base_depth, 5, strides=2,
                        padding='same', activation=tf.nn.leaky_relu),
            tfkl.Conv2D(2 * base_depth, 5, strides=1,
                        padding='same', activation=tf.nn.leaky_relu),
            tfkl.Conv2D(2 * base_depth, 5, strides=2,
                        padding='same', activation=tf.nn.leaky_relu),
            tfkl.Conv2D(4 * encoded_size, 7, strides=1,
                        padding='valid', activation=tf.nn.leaky_relu),
            # tfkl.Flatten(),
            tfkl.GlobalAveragePooling2D(),
            tfkl.Dense(tfpl.MultivariateNormalTriL.params_size(encoded_size),
                        activation=None),
            tfpl.MultivariateNormalTriL(
```

```

        encoded_size)
        # activity_regularizer=tfpl.KLDivergenceRegularizer(prior)),
    ])

    # Likelihood
    decoder = tfk.Sequential([
        tfkl.InputLayer(input_shape=[encoded_size]),
        tfkl.Reshape([encoded_size, 1, 1]),
        tfkl.Conv2DTranspose(2 * base_depth, 7, strides=1,
                             padding='valid', activation=tf.nn.leaky_relu),
        tfkl.Conv2DTranspose(2 * base_depth, 5, strides=1,
                             padding='same', activation=tf.nn.leaky_relu),
        tfkl.Conv2DTranspose(2 * base_depth, 5, strides=2,
                             padding='same', activation=tf.nn.leaky_relu),
        tfkl.Conv2DTranspose(base_depth, 5, strides=1,
                             padding='same', activation=tf.nn.leaky_relu),
        tfkl.Conv2DTranspose(base_depth, 5, strides=2,
                             padding='same', activation=tf.nn.leaky_relu),
        tfkl.Conv2DTranspose(base_depth, 5, strides=1,
                             padding='same', activation=tf.nn.leaky_relu),
        tfkl.Conv2D(filters=1, kernel_size=5, strides=1,
                    padding='same', activation=None),
        tfkl.Flatten(),
        tfpl.IndependentBernoulli(
            event_shape=input_shape,
            convert_to_tensor_fn=tfd.Bernoulli.logits),
    ])

    vae = tfk.Model(
        inputs=encoder.inputs,
        outputs=decoder(encoder.outputs[0]))

    return vae, encoder, decoder, prior

```

In [8]: `vae, encoder, decoder, prior = build_model(input_shape=train_dataset.output_shapes[0][1:])`

In [9]: `encoder.summary()`

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_6 (Conv2D)	(None, 32, 28, 28)	832
conv2d_7 (Conv2D)	(None, 32, 14, 14)	25632
conv2d_8 (Conv2D)	(None, 64, 14, 14)	51264
conv2d_9 (Conv2D)	(None, 64, 7, 7)	102464
conv2d_10 (Conv2D)	(None, 64, 1, 1)	200768

```

global_average_pooling2d_1 ( (None, 64)          0
-----
dense_1 (Dense)              (None, 152)      9880
-----
multivariate_normal_tri_l_1 ((None, 16), (None, 16)) 0
=====
Total params: 390,840
Trainable params: 390,840
Non-trainable params: 0
-----

```

In [10]: `decoder.summary()`

Model: "sequential_3"

Layer (type)	Output Shape	Param #
reshape_1 (Reshape)	(None, 16, 1, 1)	0
conv2d_transpose_6 (Conv2DTr	(None, 64, 7, 7)	50240
conv2d_transpose_7 (Conv2DTr	(None, 64, 7, 7)	102464
conv2d_transpose_8 (Conv2DTr	(None, 64, 14, 14)	102464
conv2d_transpose_9 (Conv2DTr	(None, 32, 14, 14)	51232
conv2d_transpose_10 (Conv2DT	(None, 32, 28, 28)	25632
conv2d_transpose_11 (Conv2DT	(None, 32, 28, 28)	25632
conv2d_11 (Conv2D)	(None, 1, 28, 28)	801
flatten_1 (Flatten)	(None, 784)	0
independent_bernoulli_1 (Ind	((None, 1, 28, 28), (None, 0	

```

=====
Total params: 358,465
Trainable params: 358,465
Non-trainable params: 0
-----

```

In [11]: `vae.compile(
optimizer=tf.optimizers.Adam(learning_rate=1e-3),
loss=extfp.losses.negative_log_likelihood)`

In [18]: `history = vae.fit(train_dataset, epochs=15, validation_data=eval_dataset)`

```

Epoch 1/15
235/235 [=====] - 48s 205ms/step - loss: 204.4584 - val_loss: 0.0000e+00
Epoch 2/15
235/235 [=====] - 41s 174ms/step - loss: 150.0340 - val_loss: 140.0618
Epoch 3/15

```

```

Epoch 3/15
235/235 [=====] - 41s 175ms/step - loss: 138.2131 - val_loss: 134.4889
Epoch 4/15
235/235 [=====] - 42s 177ms/step - loss: 133.0765 - val_loss: 130.2509
Epoch 5/15
235/235 [=====] - 42s 178ms/step - loss: 129.5471 - val_loss: 126.8039
Epoch 6/15
235/235 [=====] - 41s 176ms/step - loss: 127.0125 - val_loss: 125.7131
Epoch 7/15
235/235 [=====] - 42s 177ms/step - loss: 125.0497 - val_loss: 124.2229
Epoch 8/15
235/235 [=====] - 42s 177ms/step - loss: 123.3440 - val_loss: 121.7831
Epoch 9/15
235/235 [=====] - 42s 177ms/step - loss: 121.8381 - val_loss: 120.9422
Epoch 10/15
235/235 [=====] - 42s 178ms/step - loss: 120.7899 - val_loss: 120.2468
Epoch 11/15
235/235 [=====] - 42s 179ms/step - loss: 120.0660 - val_loss: 120.0946
Epoch 12/15
235/235 [=====] - 42s 178ms/step - loss: 119.4155 - val_loss: 118.6487
Epoch 13/15
235/235 [=====] - 42s 177ms/step - loss: 118.6925 - val_loss: 118.4408
Epoch 14/15
235/235 [=====] - 42s 179ms/step - loss: 118.2621 - val_loss: 117.9701
Epoch 15/15
235/235 [=====] - 42s 178ms/step - loss: 118.0233 - val_loss: 117.4856

```


```

In [19]: loss = history.history['loss']
epoch = range(1, len(loss) + 1)

plt.plot(epoch, loss)
plt.plot(epoch, history.history['val_loss'])

_ = plt.xticks(epoch)
plt.grid()

```



Visualize

```
In [0]: # We'll just examine ten random digits.
x = next(iter(eval_dataset))[0][:10]
xhat = vae(x)
assert isinstance(xhat, tfd.Distribution)
```

```
In [0]: def display_imgs(x, title=None):
    if not isinstance(x, (np.ndarray, np.generic)):
        x = np.array(x)
    n = x.shape[0]
    fig, axarr = plt.subplots(1, n, figsize=(n, 1))

    if title is not None:
        fig.suptitle(title, y=1.05)

    for i in range(n):
        axarr.flat[i].imshow(x[i].squeeze(), interpolation='none', cmap='gray')
        axarr.flat[i].axis('off')
```

```
In [22]: display_imgs(x, 'Originals')
```




```
In [23]: display_imgs(xhat.sample(), 'Decoded Random Samples')
```



```
In [24]: display_imgs(xhat.mode(), 'Decoded Modes')
```



```
In [25]: display_imgs(xhat.mean(), 'Decoded Means')
```



Check mode collapse

In [0]:

In [0]:

Anomalous Data

```
In [0]: fashion_mnist = tfds.load('fashion_mnist', split='test')
        emnist_letters = tfds.load('emnist/letters', split='test')
```

```
In [0]: def preprocess_fashion_mnist(sample):
        image = tf.cast(sample['image'], tf.float32) / 255.
        image = tf.transpose(image, (2, 0, 1))
        # image = image < tf.random.uniform(tf.shape(image))
        # image = tf.cast(image, tf.float32) - 0.5
        return image
```

```
In [0]: def preprocess_emnist(sample):
        image = tf.cast(sample['image'], tf.float32) / 255.
        image = tf.transpose(image, (2, 0, 1))
        image = tf.transpose(image, (0, 2, 1))
        # image = image < tf.random.uniform(tf.shape(image))
        # image = tf.cast(image, tf.float32) - 0.5
        return image
```

```
In [0]: fashion_mnist = (fashion_mnist
                        .map(preprocess_fashion_mnist)
                        .batch(64)
                        .prefetch(tf.data.experimental.AUTOTUNE))
```

```
In [0]: emnist_letters = (emnist_letters
                        .map(preprocess_emnist)
                        .batch(64)
                        .prefetch(tf.data.experimental.AUTOTUNE)
                        .shuffle(int(1e3)))
```

```
In [0]: def show_images(images, title=None, figsize=(4, 4)):
        if isinstance(images, tf.Tensor):
            images = images.numpy()
            images = images.squeeze()
        fig, axarr = plt.subplots(figsize=figsize, ncols=figsize[0], nrows=figsize[1])
        for ax, img in zip(axarr.flatten(), images):
            ax.imshow(img, cmap=plt.cm.gray)
```




```
ax.imshow(img, cmap=plt.cm.gray)
ax.axis('off')
if title is not None:
    fig.suptitle(title, y=1)
```


```
In [40]: x_fashion = next(tfds.as_numpy(fashion_mnist))
x_fashion_dist = vae(x_fashion)

show_images(x_fashion, title='Fashion-MNIST')
show_images(x_fashion_dist.sample(), title='Reconstructed Fashion-MNIST Sample')
```

Fashion-MNIST




Reconstructed Fashion-MNIST Sample




```
In [41]: x_letters = next(tfds.as_numpy(emnist_letters))
x_letters_dist = vae(x_letters)

show_images(x_letters, title='EMNIST/letters')
show_images(x_letters_dist.sample(), title='Reconstructed EMNIST/letters Sample')
```

EMNIST/letters



Reconstructed EMNIST/letters Sample



In [0]:

In [0]:

```
class GenerativeModel(tfk.Model):  
  
    def __init__(self, prior, likelihood):  
        super(GenerativeModel, self).__init__(name='generative_model')  
        self.prior = prior  
        self.likelihood = likelihood
```